

Epilogue: The Intelligence Economy, Re-assembled

One argument in fourteen movements

THE INTELLIGENCE ECONOMY · A WISDOM HILL RESEARCH SERIES
Wisdom Hill Research 2026-10-08

Table of contents

Preface: What This Document Is 2

1. The Economy in One Sentence 2

2. The Master Law: The Bottleneck Migrates, and Rent Follows It 3

3. Verification: The Physics of the Demand Side 4

4. The Architecture: Squeeze, Mesh, Substrate 6

5. The Consumer Mirror, and a Second Front Still Being Sown 7

6. Reading the Shadow: Two Demand Functions, Three Infrastructures 8

7. The Inversion: Why the Cycle Turns While Demand Still Rises 10

8. The Endgame and the Closed Loop 11

9. The Consolidated Ledger: Where Value Pools, and Where It Leaks 12

10. What the Series Does Not Know 13

Conclusion: Two Terminal Scarcities — Power Below, Context Above 14

Preface: What This Document Is

This series opened with an anomaly — a revenue curve without precedent in the history of commercial software — and spent fourteen reports testing whether it marked an accounting artifact or the formation of a new economic system. This closing report does not add new evidence. Its job is to show that the fourteen reports, written across two parts, in fact constitute a single argument: about what the intelligence economy is, what law governs the movement of value within it, and where that value comes to rest.

The reader who has followed the series will find no report-by-report digest in what follows — the reports are reassembled around the argument's own structure, which in Part Two happens to follow the reports' own numbering, because that part's logic is sequential; the reader who has not can treat this Epilogue as the argument in compressed form, with pointers into the reports where each claim is developed, quantified, and — this being the series' standing discipline — furnished with the conditions under which it would be falsified. Nothing below is sourced from outside the series. Where numbers appear, they are the series' own central estimates, and they inherit all the caveats stated at their origin.

One sentence of orientation. Part One (Reports 1–10) derived the intelligence economy's two demand functions — enterprise tokens and consumer media compute — from the ground up, along with the software architecture forming around them. Part Two (Reports 11–14) read those demand functions onto the physical capital stock and asked who profits from provisioning it. The Epilogue's task is to show that the two parts obey one law, and that the law has an endpoint.

1. The Economy in One Sentence

The intelligence economy exists because the unit of sale changed.

Chatbot-era AI sold answers, and answers were priced against software: the seat, at twenty dollars a month. Agentic AI sells completed tasks, and a completed task is priced against the only alternative way of obtaining it — the fully loaded cost of human labor. Between those two reference prices lie roughly two orders of magnitude, and the revenue anomaly that triggered this series is the sound of AI pricing migrating from the first anchor to the second (Report 1). Everything else in the series is downstream of this single repricing. It explains why the marginal unit of consumption became the token rather than the user; why a single developer running autonomous coding agents can consume more compute in a month than a hundred consumer subscribers; why enterprise software built to charge per human seat found itself structurally mispriced for customers that are not human (Report 4); and why the demand now propagating into semiconductors and datacenters is denominated in labor budgets, not IT budgets — the largest addressable pool ever attached to a computing workload (Report 8).

The repricing was made possible by an architectural fact the series established early and used everywhere: the agent is not a model. It is a probabilistic model wrapped in a deterministic harness — the loop, the tools, the sandbox, the context management — and it was the harness, co-evolving with the model, that carried task horizons from seconds to hours and converted

a conversational novelty into a labor substitute (Report 1). That duality of probabilistic and deterministic computation never leaves the series. It reappears as the distinction between probabilistic composition and deterministic contracts in the protocol mesh (Report 5); as the two-tier format regime of the office layer, where fluent prose and exact spreadsheets demand different guarantees (Report 7); and finally, in physical form, as the two silicon rails of Report 11 — the GPU/HBM complex that runs the model and the CPU/DRAM complex that runs the harness. The intelligence economy is not a story about intelligence alone; it is a story about the scaffolding that makes intelligence reliable enough to sell by the task.

But an input priced against labor and consumed at machine speed poses a question that a \$20 seat never did: what, exactly, limits it? The series' answer — developed empirically in Part One and structurally in Part Two — is that the limits are real, they are specific, and they move. That observation is the series' master law.

2. The Master Law: The Bottleneck Migrates, and Rent Follows It

Every report in this series, whatever its subject, performs the same analytical operation: it locates the binding constraint of its layer, observes that abundance elsewhere has made that constraint newly decisive, and concludes that economic rent is migrating toward whoever owns it. Stated once, at the level of the whole series, the law reads: **in a production system, rent accrues to the binding constraint — provided that constraint is ownable, excludable, and not quickly expanded, substituted, or regulated back toward cost; when technology dissolves one constraint, the rent does not disappear — it migrates to the next one that can be held; and agentic AI is a machine for dissolving constraints in a particular order.** Report 13 states the corollary explicitly for hardware — scarcity is where pricing power lives, and when the bottleneck migrates, the rent migrates with it — but the mechanism is the same one that organizes the software reports.

Trace the migrations the series documented, in order of appearance. Agents made *execution* hyper-abundant, so the constraint in knowledge production moved upstream to the *idea* and downstream to *verification* and *adoption* — the four-stage pipeline of Report 2, whose central finding is that flooding idea-first industries with competent execution is, in the near term and where attention is capped, largely zero-sum business stealing rather than aggregate market expansion, while the industries actually bound by execution had not yet been reached. Within coding, automated verification dissolved the human-review constraint that the complementarity literature had measured, so the constraint moved to *token efficiency per verified unit of output* — and margin began migrating to the routing, measurement, and governance layer that controls it (Report 3). Model capability commoditized across vendors, so the constraint moved first to the *harness* and then to *accumulated context* — the one asset that does not transfer when a customer switches providers (Report 1). Protocols dissolved the N^2 integration tax that had been SaaS's deepest moat, so the constraint moved to the *control plane* — and every incumbent, having opened its protocols, promptly rebuilt closure at the control and price layers (Reports 4 and 5). Open table formats commoditized data storage, so the constraint moved to *semantics* — the metric definitions, ontologies, and context layers that tell an agent what the data means

(Report 6) — and in the office estate, file formats lost their moat function to the *organizational context graph* and the knowledge network compiled from it (Report 7).

Then the law crossed from software into hardware, where Part Two gave it its sharpest formulation. Fab capacity is a self-resolving bottleneck: expensive and slow to build, but once built it produces output permanently and cumulatively. Power, grid access, and permitted sites are self-intensifying bottlenecks: they must be re-won every year against depleting inputs whose marginal cost rises as the best locations are consumed. Effective AI supply is the minimum of the two flows, and when the flows cross, the rent now earned in advanced memory and accelerators migrates downstream — toward contracted power, energized sites, and the operators who secured them before scarcity repriced them (Report 13). Even capital obeys the law: as hyperscaler funding shifted from internal cash flow to external markets, the binding constraint moved from the income statement to the balance sheet, and Report 14 found that within the investment-grade cohort the *price* of capital had become second-order while the *quantity* raisable without losing credit standing had become the constraint that ranks the field.

The law’s investment translation is uncomfortable and the series never softened it: today’s rent is evidence about the present, not a forecast. The extraordinary margins currently earned at any layer — model inference, SaaS seats, HBM — are the rent attached to today’s bottleneck, and the series’ consistent method was to ask not “who earns rent now?” but “where does the constraint go next?” That question, asked layer by layer, produced nearly every non-consensus conclusion in the fourteen reports.

3. Verification: The Physics of the Demand Side

If constraint migration is the series’ dynamics, verification is its physics. The single most load-bearing analytical device in Part One — introduced with code in Report 1, generalized in Report 2’s pipeline, formalized as the verification-times-harness framework of Report 8, and extended to media in Report 10 — is the observation that **the structure of AI demand is determined by who, or what, judges the output.**

Where verification is deterministic and machine-executed, agent work escapes the human attention budget entirely. A compiler and a test suite verify in minutes what an agent produced over hours, so autonomous retry loops can run unsupervised, token amplification per task explodes, and the domain industrializes first. This is why code fell first — not because programming is easy for language models, but because software is the one domain that ships with its own verifier (Reports 1 and 3). Where verification is human and serial, demand is throttled at the speed of reading: thirty minutes of qualitative agent work can produce a document that consumes a full day of careful review, which is why Qualitative Text Review in Report 8’s model remains the premium workload by compute weight yet capped in volume — every token it emits must ultimately pass through a human reading at two to three hundred words per minute. And where the verifier is neither a machine nor a reader but a *market mechanism*, demand explodes again by a different route: performance-advertising creative opens the window as the structural volume engine of generative media — neck and neck with the creator economy in 2026 delivered minutes — precisely because its final verifier is

the platform's A/B auction, a machine judge that evaluates thousands of variants without any human watching them. The creator economy then overtakes it on both measures and carries roughly 70% of video compute by 2028, but for the opposite reason: not a scalable verifier, but rising adoption and deeper exploration per delivered minute, bounded by the same human review bandwidth that caps qualitative work in the enterprise model (Report 10). The human visual cortex, verifying a video in half a second, sits between these poles — fast enough to consume abundance, unable to certify work.

Verification also explains the series' answer to the productivity paradox, which is worth restating because it is the hinge between the revenue anomaly and everything else. Agentic capability pooled where verification was native and deployment frictionless — code and content — which are precisely the idea-first domains where cheaper execution cannot expand a demand-fixed market. The execution-bound work of general industry, where the same capability would meet its true constraint, sat behind an enterprise-software gateway that agents could not yet inhabit. The paradox, Report 2 concluded, is the product of a topology and a misallocation: agents under the streetlight. And the constructive corollary followed from the same framework: the one domain where the logic inverts is data analytics, where ideas need not be originated but can be *found*, where agentic search industrializes discovery — and where the binding constraint therefore relocates once more, from labor to *judgment*, the interpretive discipline of distinguishing genuine signal from the spurious findings that scaled search mass-produces.

Two quantitative consequences of the verification framework organize the demand side of the series. First, the enterprise demand function: B2B LLM consumption growing from roughly 20,100 trillion tokens in 2026 to 38,100 trillion in 2028 — a 1.9-fold expansion — with software engineering's share declining from 82% to 57%. The shape of that decline is the finding, not its size: 2026 is the year agentic coding's growth peaks, after which coding decelerates to 1.18× and 1.11× annual growth while every other function grows three- to ten-fold. Coding does not retreat; it stops being the growth engine, and of the roughly 17,900 trillion tokens the seven functions add over the window, the six post-coding functions contribute 12,800 trillion. That post-coding growth is led by quantitative analysis and data science, which quadruples to roughly 8,400 trillion tokens — the one non-coding function that pairs a machine-checkable feedback loop with an autonomous-loop harness, and which therefore diffuses immediately through tool-embedded agents rather than waiting on the enterprise semantic layer, then accelerates further as that layer is built. Workflow automation grows eight-fold to roughly 3,200 trillion on the semantic layer alone, and qualitative research reaches roughly 3,000 trillion but grows only 2.9× — the ordering of growth rates tracking the verification mechanisms almost exactly (Report 8). Second, the consumer demand function: the commercial generative-media track growing roughly eight-fold from approximately 403 million to 3,392 million GPU-hours over the same window, on a different unit of account, a different verification regime, and a monetization pool capped by attention and advertising rather than wages — a cap applied throughout as an economic consistency screen on the volume assumptions, and illustrated by the one major exit of the period: OpenAI's Sora, whose consumer app was discontinued in April 2026 with the API to follow in September, and whose inference spend one outside estimate placed as high as \$15 million a day against negligible revenue, a figure never disclosed and never officially assigned as the cause, but consistent with the constraint the model imposes

(Report 10). Two demand functions, sorted by their verifiers; the whole of Part Two rests on reading both correctly.

4. The Architecture: Squeeze, Mesh, Substrate

Between the demand functions and the silicon stands the software industry, and the series devoted its longest arc — Reports 4 through 7 — to the architecture forming there. The arc has three movements: what is being squeezed, what is replacing it, and what it all stands on.

The squeeze. The SaaSocalypse repricing — roughly \$1.8 trillion of losses against nearly \$1 trillion of gains across the thirty largest software companies in six months — was directionally right about the pressure and wrong about the mechanism (Report 4). Incumbents are not refusing agentic AI; every major vendor adopted the open protocols. What they built instead is openness at the protocol layer with closure rebuilt at the control and price layers: *N* competing tollgates, each declaring itself the registry, the meter, and the trust fabric. This defense collides with a new economic fact the series called the toll problem — for the first time in software history, the marginal user reads the price list. Agents touch software thousands of times a day and can be programmed to minimize cost per task in real time, so usage pricing set at human-era rates suppresses the very machine-scale traffic that consumption models require. The result is a squeeze theorem, not an apocalypse: hold the toll line and lose share, open to machine-scale pricing and lose margin — either branch impairs enterprise value, which is why the market de-rated franchises whose quarterly results continued to beat. The dividing line in the repricing was the *denomination of revenue*: companies paid per machine actor re-rated upward; companies paid per human seat de-rated, regardless of the quality of their AI announcements.

The mesh. The gridlock is itself the strongest force feeding the alternative. Report 5 traced a single causal chain — protocols, architecture, control, pricing, money — in which standardized interaction fabrics (API, MCP, A2A, and the emerging commerce layer) collapse the integration tax; enterprise software reorganizes into a distributed operating system in which applications are system calls, agents are processes, and routers are schedulers; power concentrates in the control plane, contested at four nested scopes from the department to the economy, with the industry scope — sector semantic authority — the most underpriced position in the stack; pricing resets from seats to settlements, transmitting inference COGS into software gross margins and ending the 80%-plus gross-margin identity of the SaaS era; and settlement itself migrates to machine-native dollars, because agents cannot hold bank accounts and sub-dollar machine payments are structurally impossible on card rails. The mesh is also the redemption of a twenty-year-old failure: the service-oriented architecture of 1999–2010 lacked exactly one component, a universal semantic interpreter, and large language models supplied it — in probabilistic form, which is why reliability engineering moved from compile-time contracts to runtime harnesses, and why the harness became a locus of differentiation valuable enough that the market paid multi-billion-dollar consideration for Manus, a harness company owning no frontier model (Reports 5 and 7).

The substrate. Beneath the mesh, the data layer's fifteen-year 2x2 organization — structured versus unstructured, transactional versus analytical — collapsed along both axes under pressure from its first machine-speed customer, and value migrated upward into two new strata: the context layer that tells agents what data means, and the analytical reasoning layer where prediction, causal inference, and simulation become high-frequency agent workloads rather than scarce human projects (Report 6). Semantic lock-in emerged as the successor to data lock-in. And at the top of the substrate sits the office estate — the only horizontal software category, coextensive with the entire knowledge workforce — where the same migration takes its most consequential form: file formats are strategically spent, integration labor across document, spreadsheet, and presentation transfers to agents, and the endgame is the enterprise knowledge network, the organization's knowledge as a curated, machine-executable estate from which documents, models, and decks are compiled on demand (Report 7). Across all three movements, the same asset kept surfacing under different names — accumulated context, the semantic layer, the context graph, knowledge capital — and the series' claim is that these are one thing: the durable, non-transferable residue that remains after models, harnesses, and formats have all commoditized.

5. The Consumer Mirror, and a Second Front Still Being Sown

The consumer economy, examined in Reports 9 and 10, is not an appendix to the enterprise story; it is its controlled experiment. The enterprise reports argued that durable value forms at the control plane rather than the model layer. The consumer market validated the thesis in compressed time and at a different altitude: the consumer control plane formed at the operating system's action-brokerage layer — Android's AppFunctions shipped, Apple's App Intents declared — exactly as the enterprise control plane formed at the tool and data gateway rather than the model (Report 9). The purest expression of the logic is the period's strangest transaction: Apple licensing frontier intelligence for a reported billion dollars a year while retaining context, orchestration, and the action registry — the model treated as a component, like a display panel; the control plane treated as the franchise.

The consumer contest is really three: intent capture, action execution, and context ownership, each with its own competitive logic. The historic triumvirate of Google, Amazon, and Meta coexisted because each owned a different stage of the commerce funnel; agentic AI collapses the funnel into a single conversation, forcing each player to internalize the stages it does not own while defending its own. The period's evidence ran in one direction: execution is harder than intent. OpenAI launched checkout inside the conversation and retired it within two months; Amazon, which owns catalog ground truth, payments, and fulfillment, moved up the funnel more credibly than any intent player moved down. The base case is accordingly not a winner but a layered partition — intent distributed, execution concentrated, and tolls collected at the OS duopoly's action registries (Report 9).

Report 10 then opened the front the token lens cannot see, and closed Part One with the series' widest frame: the two axes of context conquest. On the first axis — digital work context — the war is in its harvest phase: text-centric, verification-rich, monetized against labor, measured token by token. On the second — physical life context — it is in its sowing phase: generative

media is the first monetizable beachhead, world models are the training infrastructure, and wearables are the sensor network, monetized against attention and commerce today and physical labor tomorrow.

6. Reading the Shadow: Two Demand Functions, Three Infrastructures

Part Two opens where the demand reports end: with a conversion problem. Token demand, Report 8 observed, is the shadow the enterprise sequence casts on the infrastructure market; GPU-hours are the shadow of the second front. Report 11's mandate was to read both shadows onto silicon, and its central methodological claim is that the conversion cannot be performed onto a single hardware column. AI demand lands on three infrastructures, not one: the GPU/HBM complex that executes probabilistic inference, the CPU/DRAM complex that executes the deterministic harness and orchestration around it, and the Storage/NAND complex that absorbs the data agents generate and then re-consume through the analytical flywheel. A one-column model reads the shadow onto the wrong wall.

The quantified core: the combined AI inference fleet grows from roughly 1,203,000 H100-equivalents in 2026 to roughly 3,364,000 in 2028 — a 2.8-fold expansion — of which the B2B LLM track contributes 1,095,000 rising to 2,420,000 and the generative-media track 108,000 rising to 944,000. At a reference 1.35 kW per unit, that corresponds to on the order of 1.6 GW of facility load in 2026 and 4.5 GW in 2028. These are requirement figures for twelve core workloads — Report 8's seven business functions and Report 10's five media segments, treated as peers rather than as a token core with a media appendix — plus the image and audio tracks that serve the same media customers; they are a floor on modeled demand, not a forecast of accelerator sales, and because their scope is narrower than that of the installed base examined below, which also carries training, predictive-ML, consumer serving, and uncommissioned capacity, the series deliberately never states them as a proportion of it.

The decomposition matters more than the total, in two ways. The first is physics. LLM inference is dominated by autoregressive decode: memory-bandwidth-bound, gated by HBM capacity, stacking, and interconnect. Diffusion media generation is massive-batch dense computation: in the common case compute-bound, gated by raw FLOPS, power delivery, and cooling. The series named the industry's growing practice of provisioning them separately the Great Decoupling — and was careful that the silicon buckets are not the demand tracks, since the batch-tolerant, budget-routed business functions are Token Track demand that lands on dense-compute silicon. On that cut, the HBM-bandwidth-favored core holds roughly 72% of the 2028 fleet against 28% for the dense-compute-favored segment, the latter expanding 8.8-fold over the window against 2.2-fold for the former, so its share rises from roughly 9% to roughly 28%. The Decoupling is thus a claim about increments, not convergence: the tracks do not approach parity inside the window, and dense compute's rising share of *marginal* capacity is what carries the consequences — for the HBM-versus-logic mix of semiconductor demand, for power-density planning, and for which class of silicon supplier captures the 2028 margin.

The second is the identity of the giants, and here the decomposition overturns the intuition most capacity planning is built on. Two workloads dominate, on two different silicon tracks: software engineering at 934,000 H100-equivalents rising to 1,590,000, and — once the media track is opened into its segments — the monetized creator economy at 61,600 rising to 684,200, the second-largest individual workload in the model by 2028, while professional media and entertainment, the segment a casual reading would expect to dominate, requires 17,000. Together the coder and the creator account for roughly two-thirds of all AI inference silicon in the model. And the flow diverges from the stock: of the roughly 2,161,000 accelerators added between 2026 and 2028, generative media contributes about 836,000 against software engineering's 656,000. The largest workload and the largest source of *new* demand are not the same function — which is the single most consequential fact in the series' infrastructure arc for anyone provisioning capacity rather than describing it.

Two disciplines from this conversion deserve permanent residence in any reader's toolkit, because violating either produces large errors with high confidence. The first is that **token share is not compute share**: because structured, low-wage tasks route to budget models whose per-token cost in FLOPs is five to ten times below frontier models, compute-weighted demand runs roughly 32–38% below the naive token total — approximately 13,700 trillion frontier-equivalent tokens in 2026 rising to 23,700 trillion in 2028 — so that an analysis converting nominal tokens to hardware with a single coefficient would overstate requirements by that margin. The margin is a scenario output of the model's category weights, not a measured error band, and the frontier-equivalent series rather than the nominal one is what the infrastructure volume converts (Report 8). Report 11 then showed the naive conversion errs in both directions at once: it overstates by ignoring model mix, and understates by assuming benchmark throughput, since realistic serving for the dominant interactive workloads runs at roughly half the chat-style rate. The second discipline is the generalization, the **four-distribution principle**: token share, GPU share, CPU share, and storage share are four different distributions over the same functional taxonomy, and reading any one as a proxy for the others misreads the market in proportion to their divergence — the batch-tolerant business functions holding 4.7% of 2028 nominal tokens but 0.7% of the LLM accelerator fleet, while the judgment-heavy research function does the reverse and claims a larger share of frontier-model revenue still (Report 11). Both are instances of the series' broader measurement discipline, which Part Two then applied to the installed base itself.

That application is Report 12, the series' one purely empirical exercise, and its stock-flow distinction carries the strategic finding: in 2025 the accelerator installed base tripled, from 6.6 to 20.2 million H100-equivalents, and within that torrent Nvidia captured 79% of new dollar additions while *losing* share of the cumulative stock on both compute and cost metrics — eroded from two directions, hyperscaler in-house silicon and Chinese sovereign substitution, neither of which responds to merchant-market competition because neither is a merchant-market decision. The wedge behind the first force is the decade's dominant capex fact: vertically integrated silicon delivering compute at a 2.4–5.7× per-unit cost advantage, tens of billions of dollars of avoided spend, and a 6.8-point one-year relative share swing between the integrated and merchant-dependent hyperscalers. China is the mirror image — the world's most expensive compute, carrying a rough \$22.5 billion export-control premium, substituting toward domestic supply at 85% of marginal additions regardless. Stock versus flow, here, is

not an accounting nicety; it is the difference between the market the headlines describe and the market that exists.

7. The Inversion: Why the Cycle Turns While Demand Still Rises

Report 13 is the series' theoretical center of gravity for Part Two, and its opening claim is deliberately jarring: the consensus that celebrates semiconductor suppliers as the AI buildout's structural winners and frets over the hyperscalers' compressing free cash flow accurately describes the present — and is therefore a poor guide to the future. Three frameworks carry the argument.

First, derived demand is a derivative. Gross chip demand comprises net capacity additions plus replacement of the installed base; the expansion term is purchased in proportion to the *rate of change* of AI capacity, not its level. Because adoption follows an S-curve, that expansion flow peaks near the inflection point — while headline AI demand is still setting records — and declines thereafter even as end demand grows, though the growing installed base raises the replacement floor beneath it. The industry does not need demand to fall in order to suffer; it only needs demand to decelerate, and the amplification chain running from end demand through capex to equipment orders stacks derivatives on derivatives. Second, bottleneck migration: the self-resolving constraint (fabs) crosses the self-intensifying one (power, grid, sites), and rent migrates downstream to the assets a chip glut cannot cheapen. Third, capital dynamics: cash-rich memory oligopolists face a prisoner's dilemma in which pre-emptive expansion is dominant precisely because no player is financing-constrained, while cash-consuming datacenter operators fight a war of attrition in which exit proceeds in order of capital cost — a war that concentrates the industry around its strongest balance sheets, which then inherit the scarce asset at the moment it becomes the bottleneck. The memory case study binds all three hardest, sharpened by two structural facts: a state-backed fourth player, CXMT, operating under a soft budget constraint, which forecloses the tacit supply discipline that rescued the industry in 2023, and a lead-time asymmetry that makes incumbent expansion threats credible enough to pull forward industry-wide commitment. Three supply waves converge on a 2027–28 window.

The report's honesty about what it cannot know is as load-bearing as its thesis. The economy's position on the adoption S-curve is unobservable in real time, so the peak cannot be timed — only monitored, through a dashboard of leading indicators and a trigger-based rotation design rather than a forecast. And the refined thesis that survived the report's own stress tests is fully conditional: a correction of elevated probability and amplified depth, propagating from the legacy tier toward the premium tier, with the survivors' market power durable *if and only if* the binding constraint has by then migrated to power and sites. Every clause conditional; every condition observable. That construction — a structural claim disciplined into a monitoring program — is the series' method applied to its highest-stakes call.

8. The Endgame and the Closed Loop

Report 13 specified the attrition war's exit rule but measured neither of its parameters. Report 14 measured both, and in doing so closed the loop the series opened in its Prologue. The ante per unit of capacity is the fully-loaded annual cost of operating one H100-equivalent, constructed owner by owner from a user-cost-of-capital identity in which silicon procurement sets the acquisition price, capital procurement sets the financing rate, and technology cadence sets depreciation. The staying power is financing capacity — the price of capital and, more decisively, the quantity raisable without eroding the credit standing that sets the price.

The measurement produced the series' cleanest asymmetry result — clean in direction, and, under the central normalization and cost assumptions, in approximate magnitude. The silicon axis moves the fully-loaded unit cost in *multiples*: cumulative chip cost per H100-equivalent spans nearly 2× across the major hyperscalers, produced almost entirely by in-house silicon, and dilutes only to roughly 1.6× after uniform infrastructure is layered on. The capital axis, within the investment-grade cohort, moves the same cost in *basis points* — one to two percent — but rations in dollars: the 2025–26 funding-model inflection, with hyperscaler bond issuance running at four times the prior five-year average and projections of \$400 billion of borrowing in 2026, shifted the binding constraint to the balance sheet, where the cohort diverges sharply for reasons that mostly predate AI. The composite verdict named Google as the only dual-axis leader — cheapest fully-loaded compute at roughly \$5,000 per H100-equivalent per year against roughly \$8,000 at Microsoft and nearly \$10,000 at xAI and CoreWeave, combined with demonstrated access to every layer of the capital structure — and flagged, on the other side of the ledger, the cohort's least-priced risk: the counterparty circularity around OpenAI, whose solvency is referenced by equity value, contracted cloud revenue, and capacity planning alike. The marginal cost of compute lives in the neocloud tail, financed against collateral depreciating 30% a year; any demand disappointment reprices there first.

One convergence gives the endgame its significance beyond a leaderboard, and it is the loop itself. The Prologue promised that the series' two parts were causally coupled: infrastructure cost structures set the price of tokens; token prices set the economics of agents; agent economics determines how far labor substitution proceeds and therefore how much infrastructure is demanded. Report 14 supplied a time-stamped estimate of the loop's most important resistance — the cost floor of the intelligence economy is set by vertically integrated fleets, and the gap to the rest of the field is a multiple, not a margin. This floor is not a fixed constant: it drifts with chip generation, workload mix, utilization, power cost, and software efficiency, and Report 14 measured it at one point in time rather than for all time. Every downstream thesis of Part One — token-price deflation, agent unit economics, the pace at which the toll problem squeezes incumbent software, the breakeven arithmetic of substituting agents for labor — inherits that floor. The economy the series describes is not a stack of independent layers; it is a circuit, and the circuit now has a resistance that has at least been measured once.

9. The Consolidated Ledger: Where Value Pools, and Where It Leaks

Assembled across all fourteen reports, the series' value-migration map reduces to a small number of positions, each an instance of the master law. The table states each position, the constraint whose ownership sustains it, and the report where it is developed and given falsifiers; the prose that follows states the pattern.

Position	The constraint it owns	Developed in
Control planes (routing, orchestration, governance) — at token, enterprise, industry, and OS scopes	Access to and allocation of scarce agent capability	Reports 3, 5, 9
Context and semantics — accumulated agent context, semantic layers, knowledge capital	Meaning: the non-transferable residue after models and formats commoditize	Reports 1, 6, 7, 8
Verification infrastructure — test harnesses, semantic layers as machine verifiers, audit surfaces	The gate through which abundant execution becomes trusted output	Reports 2, 3, 8
Settlement and trust rails — metering, machine-native payment, agent identity	The transactional plumbing seats never needed	Reports 4, 5
Machine-denominated revenue models	Exposure to the number and activity of agents, not human seats	Report 4
Vertically integrated, strongly financed compute operators	The cost floor of compute and the staying power to keep building	Reports 12, 14
Power, sites, and the equipment chain behind them	The self-intensifying downstream bottleneck	Report 13
Human judgment and frontier ideas	The upstream constraint no execution abundance relieves	Report 2

Value leaks, correspondingly, from every position whose product was a constraint that agents dissolve: integration middleware whose product was the integration tax itself; commoditized, routable model inference; thin-workflow SaaS whose function an agent composes from primitives; seat-denominated revenue at machine-scale usage; file formats as moats; merchant-dependent compute procurement carrying a multiple-not-margin cost disadvantage; leveraged capacity financed against fast-depreciating collateral; and — conditionally, on Report 13's triggers — the memory and accelerator rents that today look most secure. The common signature of a leaking position is that it earns rent on a bottleneck whose dissolution is already funded.

Two entries in the ledger deserve emphasis because they bracket all the others. At the top of the stack, the series' longest-running conclusion is that after each successive commoditization — model, harness, format, protocol — the residual asset is always a form of *context*: what the organization knows, what its data means, what its agents have accumulated. At the bottom,

the terminal asset is *power and sites*: the one input a technology cycle cannot manufacture and a chip glut cannot cheapen. Everything in between is contested, and mostly contestable.

10. What the Series Does Not Know

A synthesis owes the reader its residual uncertainties, stated as plainly as its conclusions.

The paradox is answered structurally, not settled empirically. Report 2 resolved the productivity paradox's composition — a topology and a misallocation — and identified where the latent gains live. But Report 1's Coda posed a harder question the series could bound without closing: whether the revenue curves of 2026 are the leading edge of a productivity J-curve or the signature of a capital cycle buying expenditure rather than outcomes. The series' answer is conditional and monitorable rather than declarative: the mechanism for genuine value is identified (execution-bound work, reached through the mesh, verified through the semantic layer, interpreted with judgment), the routing obstacle is identified (the enterprise gateway, and the gridlock of Report 4), and the falsifiers are stated in each report. If measured productivity in execution-bound functions fails to respond as the gateway opens through 2027, the capital-cycle reading strengthens — and Part Two's inversion frameworks say what happens next.

Timing is structurally unknowable, and the series refused to fake it. The position of the economy on the adoption S-curve is unobservable in real time; the inversion's peak, the memory correction's onset, and the attrition war's casualties are monitorable through stated dashboards, not datable in advance. Where the series names windows — the 2027–28 convergence of supply waves — they are the intersection of observable construction and ramp schedules, not demand forecasts.

The load-bearing numbers carry declared error bars. The installed-base and cost figures underlying Part Two are median estimates with confidence intervals spanning roughly a factor of two. That uncertainty is real enough that the hardware conclusions are stated at the level of direction, rankings, and gap multiples under the central normalization and cost assumptions, rather than as point values whose exact ordering would survive every joint or platform-specific revision. The demand functions of Reports 8 and 10 are built from anchored bases with graded confidence, and their weakest links are flagged at source: the penetration paths and category compute weights on the enterprise side, and on the media side an explicitly conditional scenario whose totals may plausibly vary by several-fold, with creator adoption, delivered minutes per video, and assumed exploration depth together setting roughly 70% of the track. The fleet counts that convert both are capacity-equivalents rather than device counts — production serving mixes GPUs, TPUs, and inference ASICs, and the media coefficients are engineering assumptions calibrated to published throughput and pricing rather than hardware measurements — and they are a floor on twelve modeled workloads, never a forecast of accelerator demand. Anthropic's revenue figures, which triggered the series, remain run-rate disclosures with documented tensions against recognized revenue; the series treated them as evidence of a repricing, never as audited fact.

And the strongest conclusions have named enemies. Correlated failure across the verification stack would partially restore the human-review regime and re-throttle the coding demand function. Per-task pricing at the model layer would claw back the margin this series routed to the control plane. Regulatory action could cap the gateway rents, force open the action registries, or break the bundles. A supply-discipline surprise in memory, or a demand-side escape along the Jevons path, would defer the inversion. Each is stated, in its home report, with the observable that would confirm it. A series that claims an industrial revolution owes its readers the shape of its own refutation; this one has tried to pay that debt in every report.

Conclusion: Two Terminal Scarcities — Power Below, Context Above

Strip away the layers this series has traversed — the agents and their harnesses, the mesh and its tollgates, the semantic strata, the two demand functions, the three infrastructures, the stock and the flow, the war of attrition — and, over the investment horizon this series examines, the intelligence economy resolves into a system bounded by two scarcities that no near-term abundance of intelligence relieves.

Below, the boundary is physical: energy, and the sited, permitted, grid-connected capacity to deliver it. Every constraint beneath the agent is self-resolving on some horizon — fabs get built, packaging scales, models compress — while this one scales slowest of all: not impossible to expand, but bottlenecked by generation, transmission, and site development that move at the pace of physical infrastructure and are limited region by region, so that it intensifies as it is consumed faster than it is relieved. The series' hardware arc is, in the end, the story of rent migrating toward that floor, and of the balance sheets durable enough to be standing there when it binds.

Above, the boundary is proprietary: context, and the control planes that govern access to it. Agentic AI makes execution abundant, search cheap, and generation industrial — and in doing so it manufactures, at machine speed, the one commodity that was never scarce: plausible output. What it cannot manufacture is the material that makes output trustworthy and actionable, because that material must be true of a particular firm or a particular person rather than merely plausible. In the enterprise it is the semantic layer that fixes what a customer, an order, or a margin actually means inside one company, and the verification infrastructure that certifies which outputs can be acted on without a human reading them. In the consumer economy it is the personal context accumulated through a person's use of their devices — their history, their preferences, and, as wearables become the sensor network of the second axis, the record of their physical life — together with the action registries through which an intent becomes a transaction and the catalog, payment, and fulfillment ground truth that settles what was in fact bought. None of it can be generated; all of it is accumulated — through data governance and permissioning inside the firm, through use on the consumer side — and once accumulated it is owned by whoever controls the plane through which the activity passes. That is why the series found value forming at the control plane rather than the model layer at every level it looked: in the enterprise gateway that releases the post-coding functions, in the

settlement rails of the mesh, in the operating system's action-brokerage layer, in the platforms that own the generation model and the pool that monetizes it alike. The series' software arc is the story of rent migrating upward toward that ceiling — from execution to verification to context — until the residual claimant on the stack is whoever controls the context the stack is acting on.

An economy provisioned like power and priced against labor will be built between these two boundaries. The fourteen reports of this series have mapped the territory in between: where the constraints sit today, the order in which they dissolve, and the positions — control planes, context, verification, settlement, integrated compute, contracted power — where the rent comes to rest at each stage. Human judgment sits above that map rather than on it, choosing which frontiers are worth the power and the context in the first place; that is a question about what the economy is for, and this series has been about what it costs. The map will be wrong in places; the falsifiers say where to look. But the law that drew it has held at every layer the series examined, and it is the sentence with which the series should close. The bottleneck always moves — and value, patient and indifferent to narrative, always follows it.

This Epilogue synthesizes Reports 1–14 of The Intelligence Economy series and introduces no evidence beyond them. All figures are the series' own central estimates and inherit the confidence grades, measurement caveats, and falsification conditions stated in their originating reports.

The Intelligence Economy

© 2026 Wisdom Hill. Licensed under CC BY-NC-ND 4.0.

Quotation and summary are freely permitted with attribution.

Cite as: Wisdom Hill Research. 2026. "Epilogue: The Intelligence Economy, Reassembled."

In The Intelligence Economy. Wisdom Hill.

<https://wisdomhill.github.io/intelligence-economy/>