

Prologue: About This Series

An Anomaly, a Claim, and the Architecture of Fourteen Reports

THE INTELLIGENCE ECONOMY · A WISDOM HILL RESEARCH SERIES
Wisdom Hill Research 2026-08-24

Table of contents

The Trigger 2

The Claim 2

The Architecture: What Changed, What Is Being Built, What It Runs On 3

How to Read the Series 5

References 6

The Trigger

Every research program begins with an anomaly. Ours is a revenue curve.

Anthropic's annualized run-rate revenue — the most recent month's revenue multiplied by twelve, not audited revenue under Generally Accepted Accounting Principles (GAAP) — stood at approximately \$9 billion at the end of 2025. The company disclosed \$14 billion in its February 12 Series G announcement and \$30 billion on April 6; press reporting places the intervening March reading at \$19 billion. On May 28, in announcing a \$65 billion Series H round at a \$965 billion post-money valuation, the company confirmed that run-rate revenue had crossed \$47 billion earlier that month. More than a fivefold increase in under five months. Meritech Capital's Alex Clayton, who had examined the initial public offerings of over two hundred public software companies, told CNBC in June 2025 that he had never observed a growth rate of this kind — a judgment rendered when the run-rate stood near \$3 billion, eleven months and some fifteenfold below the May 2026 figure.

Nor does the signal rest on one company's self-reporting. On February 3, 2026, a single product release erased roughly \$285 billion of market capitalization in one trading session: the selling opened in legal- and data-services equities before spreading across software, financial services, and asset management, and the Goldman Sachs basket of U.S. software stocks fell 6% (Report 1, Section 2.3.3). And over six months, on our tabulation, the thirty largest software companies were repriced by roughly \$2.8 trillion in absolute value — \$1.8 trillion of losses against nearly \$1 trillion of gains, sorted not by AI exposure but by the denomination of revenue. Report 4 sets out that tabulation in full: the company list, the measurement window, and the sorting rule. Sellers of AI and sellers of the things AI displaces are being repriced simultaneously, in opposite directions.

An anomaly of this scale admits two explanations. Either the numbers are an artifact — of run-rate accounting, of circular financing, of a capital cycle marking itself to fantasy — or something has changed in the nature of what is being sold. This series argues, with evidence and with explicit falsification conditions, for the second explanation. The first is never dismissed; it is tested repeatedly. Report 2 confronts its strongest form directly, and Report 14 returns to it at the layer where a capital cycle would ultimately break — the balance sheets funding the buildout.

The Claim

What changed is the unit of sale.

Chatbot-era AI sold *answers*, and answers were priced the only way software knows how to price: the seat, at \$20 per month. Agentic AI sells *completed tasks* — a refactored codebase, a reconciled ledger, a drafted contract reviewed against precedent — and a completed task carries a different reference price: the fully loaded cost of the human labor that would otherwise have performed it. Between the software seat and the knowledge-work wage lie roughly two orders of magnitude. The 2025–2026 revenue explosion is, at bottom, the sound of AI pricing migrating from the first anchor to the second. The consumption data confirm the

mechanism: the marginal unit of AI consumption is no longer a user but a token, and a single developer running an autonomous coding agent consumes hundreds to thousands of dollars of compute per month — on our arithmetic, one work-referenced API user can be worth a hundred consumer seats.

When the price of a produced good collapses by orders of magnitude while its quality threshold is crossed, economies reorganize around it. This happened with mechanical power, with electricity, and with computation itself. We use the term **intelligence economy** for what is now forming: an economic system in which cognitive task-completion becomes a purchasable, meterable input — provisioned like power, priced against labor, and consumed in volumes that are already bending the physical supply chains beneath it. Calling this an industrial revolution is a strong claim, and the series holds itself to a corresponding discipline: every structural thesis it advances is accompanied by the conditions under which it would be falsified.

The Architecture: What Changed, What Is Being Built, What It Runs On

An input that is produced, consumed, and physically provisioned must be studied at all three levels — but it must also be studied in two tenses, because part of this transformation is already on the record and part of it is not. The series therefore cuts twice. The first cut is temporal: what can already be measured, against what is still forming. The second cut, applied to what is still forming, is by layer: the software architecture rising into the vacated space, and the physical capital stock that has to carry it.

Part One — The Present Tense: What Has Already Changed (Reports 1–4) establishes what is observable today. Reports 1 and 2 set out the phenomenon and its central puzzle: the capability transition from conversation to labor, and the gap between exploding output and stagnant measured productivity — a gap the series resolves into a structural claim (execution has become abundant while ideas, verification, and judgment have not) and a compositional one (agents pooled where deployment was easy, not where value was large). Reports 3 and 4 examine the beachhead and the incumbents: why code fell first and what its token economics reveal, and why the enterprise-software estate that guards the remaining work is squeezed whichever way it turns. The load-bearing evidence in this part is observed prices, observed consumption, and observed repricing; where it forecasts, it does so one step beyond data already in hand.

Part Two — The Forward Vision: The Architecture Rising in Its Place (Reports 5–10) turns from what has happened to what is being built on the vacated ground. Reports 5 through 7 map the emerging structure — the protocol mesh and the control-plane stack above it, the converging data layer and its new semantic strata beneath it, and the office layer where organizational knowledge itself is being recompiled. Reports 8 through 10 convert that structure into demand: Report 8 quantifies enterprise consumption token by token across seven business functions; Reports 9 and 10 turn to the consumer economy, where the control plane forms at the operating system's action registry and a second demand function — generative media, measured in GPU-hours rather than tokens — opens the physical-context front. These

are projections rather than observations, and each is stated with the conditions that would falsify it.

Part Three — The Physical Substrate: The Infrastructure of Intelligence (Reports 11–14) reads those demand functions onto the physical capital stock. Report 11 converts the two demand series of Part Two into hardware requirements across three distinct infrastructures — accelerators and their high-bandwidth memory (GPU/HBM), general-purpose compute and system memory (CPU/DRAM), and storage (NAND flash) — and shows why no single coefficient can do the job. Report 12 measures the accelerator installed base itself: stock versus flow, supplier and owner shares, and the vertical-integration cost wedge that is quietly re-ranking the hyperscalers. Report 13 supplies the cycle theory — why derived demand is a derivative, why bottlenecks migrate downstream from fabs toward power and sites, and why the capital dynamics of memory suppliers and datacenter operators point toward an inversion of today’s consensus. Report 14 names the inheritors: a two-axis analysis of silicon and capital procurement that ranks who acquires compute most cheaply and who can keep acquiring it longest. Part Three is grouped by layer rather than by tense, and therefore holds both a census and two forecasts: Report 12 measures the installed base on which Reports 13 and 14 build.

The three parts are not adjacent topics; they are causally coupled. Infrastructure cost structures set the price of tokens; token prices set the economics of agents; agent economics determines how far substitution proceeds and therefore how much infrastructure is demanded — a loop in which Jevons dynamics, margin migration, and capital-cycle risk all live. The series closes this loop explicitly in Report 14, and the Epilogue reassembles the whole into a single argument.

Part One — The Present Tense: What Has Already Changed

What is already on the record: the unit of sale moved, and the incumbent estate has been repriced for it.

#	Report	The question it answers
—	Prologue: About This Series (<i>this document</i>)	Why does this series exist, what does it claim, and how is it built?
1	The Agentic Inflection	What changed in 2025–2026, and who is fighting over the agentic value layer?
2	The Idea Bottleneck	Output is exploding — why is measured productivity not, and where does the dividend actually hide?
3	The Token Economics of Coding Agents	Why did code fall first, and what governs token consumption as the unmetered era matures?
4	The SaaS Apocalypse Question	Is incumbent software dying — and what is the toll problem that squeezes it either way?

Part Two — The Forward Vision: The Architecture Rising in Its Place

What is being built on the vacated ground, and how much intelligence it will consume.

#	Report	The question it answers
5	The Agentic Mesh	How does protocol standardization rewire software architecture, control, pricing, and money?
6	The Data Layer Reforged	As agents become the data layer's first customer, where does its value migrate?
7	The Agentic Office	What happens to office software, the duopoly, and the organizational knowledge they hold?
8	The Functional Anatomy of Enterprise AI Demand	Function by function, how many tokens does enterprise work consume through 2028?
9	Who Owns the Consumer Agent?	When the commerce funnel collapses into a conversation, who captures intent, execution, and context?
10	The Functional Anatomy of Consumer AI Demand	How large is generative-media demand — and what do world models and wearables open beyond it?

Part Three — The Physical Substrate: The Infrastructure of Intelligence

Where that demand lands: silicon, memory, power, and the balance sheets that procure them.

#	Report	The question it answers
11	Reading the Shadow onto Silicon	How do two demand functions land on three distinct hardware infrastructures?
12	The 2025 AI Accelerator Market	What do stock and flow reveal about supplier power and the vertical-integration wedge?
13	The Coming Inversion of the AI Hardware Cycle	Why does the hardware cycle turn while demand still sets records — and where do the bottlenecks migrate?
14	The Hyperscaler Endgame	Who inherits the inversion: who builds compute cheapest, and who can fund it longest?
—	Epilogue	What do the fourteen reports add up to — one economy, one law, two terminal scarcities?

How to Read the Series

Four instruments recur across all fourteen reports, and recognizing them will make any single report easier to read. First, **constraint migration**: each report locates the binding constraint of its layer and argues that economic rent moves with it — from execution to ideas and judgment, from models to harnesses to context, from data formats to semantics, from chips to power and sites. Second, **verification** as the physics of the demand side: what can be machine-checked scales without human limit; what must be humanly read or watched cannot, and the entire structure of AI demand sorts along that line. Third, the **control plane**: at every scope — the token budget, the enterprise, the industry, the operating system — value concentrates in the layer that governs access, routing, and settlement rather than in the commoditizing

capability beneath it. Fourth, **measurement discipline**: run-rate is not audited revenue under Generally Accepted Accounting Principles (GAAP), token share is not compute share, stock is not flow, and the series flags on every load-bearing number whether it is measured, estimated, or judged.

Readers pressed for time can take three paths, each cutting across the three parts rather than following them. The *economic* path — Reports 1, 2, 3, 8, 13 — carries the demand thesis from inflection through token economics to inversion. The *strategic* path — Reports 4, 5, 6, 7, 9, 14 — follows the competitive contests for the control planes and the data layer beneath them. The *infrastructure* path — Reports 8, 10, 11, 12, 13, 14 — runs the demand functions down to silicon and capital. All paths converge in the Epilogue.

The series begins where the anomaly begins. Report 1 establishes the two facts on which every subsequent report depends: what changed, and who is fighting over it.

References

All sources accessed August 2026. Run-rate figures are company-reported annualizations of a single month's revenue, not audited revenue under Generally Accepted Accounting Principles (GAAP).

1. Anthropic — “Anthropic raises \$30 billion in Series G funding at \$380 billion post-money valuation” — run-rate revenue of \$14 billion (Feb 12, 2026). <https://www.anthropic.com/news/anthropic-raises-30-billion-series-g-funding-380-billion-post-money-valuation>
2. Anthropic — “Anthropic expands partnership with Google and Broadcom for multiple gigawatts of next-generation compute” — run-rate revenue surpassing \$30 billion, against approximately \$9 billion at end-2025 (Apr 6, 2026). <https://www.anthropic.com/news/google-broadcom-partnership-compute>
3. Anthropic — “Anthropic raises \$65B in Series H funding at \$965B post-money valuation” — the \$65 billion raise, the \$965 billion post-money valuation, and run-rate revenue crossing \$47 billion earlier in May (May 28, 2026). <https://www.anthropic.com/news/series-h>
4. VentureBeat — “Anthropic says it hit a \$30 billion revenue run rate after ‘crazy’ 80x growth” — run-rate trajectory from January 2024 to April 2026, including the March 2026 reading of \$19 billion (May 8, 2026). <https://venturebeat.com/technology/anthropic-says-it-hit-a-30-billion-revenue-run-rate-after-crazy-80x-growth>
5. CNBC — “Anthropic: 2025 CNBC Disruptor 50” — Alex Clayton (General Partner, Meritech) on the absence of precedent across more than two hundred software IPOs, at a then-prevailing run-rate of roughly \$3 billion (Jun 10, 2025). <https://www.cnbc.com/2025/06/10/anthropic-cnbc-disruptor-50.html>
6. Bloomberg — “Anthropic AI Tool Sparks Selloff From Software to Broader Market” — roughly \$285 billion of market value erased in a single session across software, financial services, and asset management; the 6% decline in the Goldman Sachs U.S. software basket; the selloff opening in legal and data services and spreading outward (Feb 3, 2026). <https://www.bloomberg.com/news/articles/2026-02-03/legal-software-stocks-plunge-as-anthropic-releases-new-ai-tool>

7. CNBC — “Anthropic tops OpenAI as most valuable AI startup, nears \$1 trillion valuation in latest round” — independent reporting of the Series H round and the resulting standing against OpenAI (May 28, 2026). <https://www.cnbc.com/2026/05/28/anthropic-open-ai-startup-value.html>

The \$2.8 trillion six-month repricing of the thirty largest software companies is Wisdom Hill Research’s own tabulation from exchange closing prices and shares outstanding; the company list, measurement window, and sorting rule are set out in Report 4.

The Intelligence Economy

© 2026 Wisdom Hill. Licensed under CC BY-NC-ND 4.0.

Quotation and summary are freely permitted with attribution.

Cite as: Wisdom Hill Research. 2026. “Prologue: About This Series.”

In The Intelligence Economy. Wisdom Hill.

<https://wisdomhill.github.io/intelligence-economy/>