

10. The Functional Anatomy of Consumer AI Demand

Generative media economics, world models, and the conquest of physical-life context, 2026–2028

THE INTELLIGENCE ECONOMY · A WISDOM HILL RESEARCH SERIES

Wisdom Hill Research 2026-08-28

Table of contents

1. Executive Summary	2
2. Introduction	4
2.1 Scope and the meaning of “consumer”	4
2.2 Why a separate compute track	5
2.3 The integrating frame: two axes of context conquest	5
3. The Framework Extended	6
3.1 Verification at a glance	6
3.2 The exception that drives the volume: the auction as machine verifier	6
3.3 The monetization-pool cap	6
3.4 The unit of account and the shape of the compute	7
4. The Supply Side	8
4.1 The contenders — and OpenAI’s exit	8
4.2 Why video is multi-polar where LLMs were not	9
4.3 The cost collapse, and the status of the GPU-hour coefficients	9
5. The Functional Anatomy of Media Demand	10
5.1 Method: anchor, penetrate, deepen — then screen against the pool	10
5.2 Why the penetration rates differ: denominators, deciders, and forcing functions	14
5.3 Segment 1 — Performance Ad Creative: supplier-push, not customer-pull	15
5.4 Segment 2 — E-commerce SKU Video: lowest penetration of stock, highest penetration of commerce	16
5.5 Segment 3 — Creator Economy / UGC: the depth-driven leader from 2027	17
5.6 Segment 4 — Professional M/E: the highest headline penetration and the Netflix counter-intuition	18
5.7 Segment 5 — Corporate Internal: the structural laggard, honestly sized	20
5.8 The three compounding curves behind every path	20
5.9 The platform synthesis: three pool-owners, one two-axis competitor	21
5.10 The image and audio tracks	21
5.11 Totals and the fleet equivalent	22
6. Beyond Content	23
6.1 From video models to world models: Google’s integrated path	23
6.2 Glasses and microphones: the interface war moves onto the body	23
6.3 The two-axis synthesis	24
7. Risks and Sensitivities	25
8. Tracking Indicators	26
References	27

1. Executive Summary

Report 8 of this series mapped enterprise AI demand across seven business functions, all sharing a common unit of account — the LLM token — and a common economic logic: agents priced against the labor budgets of knowledge work. This report maps the demand function that the token lens cannot see. Generative media and its successors run on a different unit of account (GPU-hours of visual and audio generation, not convertible to tokens), a different verification regime (the human visual cortex, or the advertising auction, rather than the compiler), a different model architecture (compute-bound parallel diffusion rather than memory-bandwidth-bound autoregressive decode), a largely disjoint supplier set (Google, ByteDance, Alibaba, Kuaishou, and independents; Anthropic is absent), and a different monetization logic, centered on consumer attention and commerce rather than primarily on wage budgets. Because the supplier and measurement ecosystems are disjoint, LLM-derived usage statistics structurally undercount this entire economy.

Three headline results organize the analysis. First, **the commercial media track modeled here — five video segments plus the image and audio tracks that serve them — grows roughly eight-fold over the forecast window, from approximately 403 million GPU-hours in 2026 to 3,392 million in 2028** — an equivalent fleet expanding from roughly 108,000 to 944,000 H100-equivalent inference accelerators at segment-specific utilization rates — driven by generated video volume rising eight-fold from approximately 922 million to 7,835 million minutes per year, plus image generation scaling from 50 to 150 billion units and voice and music from 0.5 to 2.5 billion hours. Every video figure is built forward from a stated reference anchor in the existing online-video economy — advertiser accounts, marketplace listings, creator populations, professional production flows, and organizational counts — with explicit penetration paths. For four of the five segments the figure is expressed first in *delivered* minutes, the minutes that actually reach an audience, and compute is derived from them through an explicit two-pass structure: a segment-specific number of exploration passes at workhorse quality, plus one-and-a-half to two finishing passes at the quality the segment actually delivers in — a light upgrade over workhorse where the consistency requirement is low, full premium tier for professional footage. Professional media and entertainment is modeled differently, because almost none of what it generates reaches the screen: its exploratory workload is estimated per adopting project and its delivered footage separately. The distinction matters because the discarded drafts that dominate volume are not rendered at the quality of the frame that survives, and the gap between the two is widest precisely in the segments that explore hardest. Because production throughput and adoption are only partially observable, the totals should be read as an order-of-magnitude conditional scenario under the stated assumptions, plausibly varying by several-fold, rather than as a point forecast. The trajectory is far faster in relative terms than the enterprise LLM track: an eight-fold expansion in GPU-hours against the roughly two-fold growth of compute-weighted B2B token demand. These are five commercially monetized segments, not the media economy as a whole — personal and non-commercial generation, larger than all of them on a units basis, is outside the estimate for the reasons given in Section 2.1.

Second, **the structure of the demand inverts the intuitions imported from both the LLM market and the traditional media industry**. The structural volume engine at the opening of the window — running neck and neck with the creator economy in 2026 delivered minutes —

is performance-advertising creative, the media workload whose final verifier is not a human but the platform's A/B auction: a scalable statistical verifier that lets variant volume escape the human-review bottleneck, as the verification framework of the companion enterprise report anticipates. The fastest-growing segment, and the one that dominates compute by 2028, is the creator economy — YouTube channels, short-form operators, video podcasters — whose demand is driven not by a rising head-count but by rising adoption and by more delivered AI minutes in each video, amplified by an exploration depth that is high in every year rather than rising, and which reaches roughly 70% of video GPU-hours by 2028. Professional media and entertainment inverts in the opposite direction, and more sharply than any other segment: it contributes under one tenth of one percent of delivered minutes yet roughly 2.5% of video GPU-hours — an over-representation of about thirtyfold, driven by premium intensity and very deep idea exploration, but still not the compute-dominant segment that a “Netflix drives the market” intuition implies. The reason is structural: studio use of generative video is overwhelmingly *exploratory* — tens of thousands of short candidate clips per project, spread across hundreds of shots and sequences, generated at low quality and almost never reaching the screen — while what does reach the screen is largely AI-assisted post-production on conventionally shot footage. Each segment is built forward from a stated reference anchor of varying evidentiary strength (Meta's ~11 million advertisers, YouTube's ~500-hours-per-minute upload firehose and ~15 million regular creators, the ~2 billion marketplace listings, ~60,000 annual professional productions, and ~4 million organizations) with an explicit penetration logic. And the supply side is genuinely multi-polar where the LLM market was not — Google's Gemini Omni Flash currently holds the top leaderboard position while Chinese vendors (ByteDance, Alibaba, Kuaishou) fill most of the remaining top tier, reflecting the short-video giants' possession of unusually large proprietary video datasets and the distribution channels that monetize the output; OpenAI, meanwhile, has exited the consumer-facing Sora app and is completing its withdrawal from video generation, with the API scheduled to shut down on September 24, 2026.

Third, **the content market is the near-term P&L, but it is not the endgame.** The same model lineage extends along a technical continuum from video generation to interactive world models — Google's Genie, already training embodied agents in simulated environments — and the interface war is moving onto the body, with Meta's AI glasses reaching seven million-plus units on the sales measure reported for 2025, alongside IDC's separate forecast of 13.6 million display-less smart-glasses shipments in 2026. The two figures come from different series and are not directly comparable — Section 6.2 sets out the caveats — but both point the same way. Read through the two-axis context framework that closes this report, generative media is the opening battle of a second front: where the enterprise agent conquered *digital work context* and monetizes against labor, multimodal systems, wearables, and world models are the machinery for acquiring *physical life context*, monetizing against attention and commerce today and physical labor tomorrow. For the attention-funded segments that carry most of the track's volume, the binding constraint is not capacity to generate but capacity to monetize — a constraint applied throughout this report as an economic consistency screen on the volume assumptions, and illustrated by OpenAI's exit from consumer video: the Sora app was discontinued in April 2026 with the API scheduled to follow in September, a retreat consistent with — though never officially attributed to — inference economics that one outside estimate placed as high as \$15 million per day against negligible revenue. One structural pattern shapes

every scenario: Google currently fields the most visibly integrated stack across both axes — consumer distribution, generative video, world models, robotics models, and wearable partnerships — a breadth no single competitor presently matches within one ecosystem.

2. Introduction

2.1 Scope and the meaning of “consumer”

This report is the consumer-side companion to *The Functional Anatomy of Enterprise AI Demand* (Report 8), which partitioned B2B work into seven token-denominated functions and deliberately excluded an eighth — Generative Media — from its quantitative core. This report develops that eighth function in full and extends it to the frontier that follows it.

A definitional note first, because the word “consumer” carries two meanings here and the distinction disciplines the model. Much of the generated-media volume analyzed below is *purchased* by businesses — advertising agencies, e-commerce sellers, corporate learning departments. What makes the demand “consumer” in the economically relevant sense is where the media is distributed and monetized, not that every production budget originates directly in consumer attention. The relationship runs in three degrees. Performance advertising and creator media are directly attention-funded, drawing on advertising, subscription, and attention dollars that originate with consumers rather than on the wage and labor budgets that fund the enterprise agent economy. E-commerce and professional media are one step removed: they are funded from conversion economics and production budgets, but both exist to move product or programming to consumers, and their volume tracks consumer demand. Corporate internal video is the bridge case, included because it runs on the same generative-media production stack and competes for the same accelerators, though its funding is an enterprise operating budget and its constraint is organizational habit rather than monetization. The first degree is what gives the model its discipline. This distinction is not cosmetic. The attention pool supplies the single most important constraint in the model — the monetization-pool cap of Section 3.3, which binds the first-degree segments directly and screens the rest — and it explains why the demand function behaves differently from the token economy: enterprise agents are bought against labor costs that total tens of trillions of dollars globally, while generated content competes for a fixed pool of audience attention and the advertising and subscription revenue attached to it.

Scope of the quantitative estimate. The demand modeled in Section 5 is not the whole of generative media. Five video segments are modeled quantitatively, supplemented by separate image and audio tracks that serve the same customers; the largest source of generation today is not among them.

We exclude personal and non-commercial creation entirely — images, music, and video generated by individuals for learning, entertainment, self-expression, and social sharing. On any current measure of units generated this is the largest block in the market, and free tiers make it the least monetized. It is excluded because no commercial base population exists against which to anchor it, and because the monetization cap of Section 3.3 — the discipline that makes the rest of the model defensible — does not apply to generation that nobody is paying for.

Among commercially motivated production we model five segments: self-serve performance advertising, e-commerce product media, the monetized creator economy, professional media and entertainment, and corporate internal video. Others are absent. Agency- and brand-side exploratory generation is discussed qualitatively in Section 5.3 but not counted. Day-to-day broadcast operations — news, sports graphics, weather, educational and religious programming, where AI-generated explanatory footage is growing — are outside the professional production flow this model counts. Game asset production, architectural and product visualization, news and publishing media, and education content production are all plausibly material and all omitted.

The figures in Section 5 should therefore be read as a floor on the five segments named, not as an estimate of generative media demand. Section 6, which extends the analysis to world models and the physical-context frontier, is not bounded by that scope: it is where the consumer side of this report's title does most of its work.

2.2 Why a separate compute track

Three properties compel separate treatment rather than a footnote to the token model. The **unit of account** differs: visual generation is denominated in spacetime patches and audio frames, not language tokens, and no defensible exchange rate converts one into the other; we therefore measure in GPU-hours. The **model architecture** differs: as a general tendency, diffusion-class generation is a compute-bound, massively parallel batch workload, against the memory-bandwidth-bound autoregressive decode of LLMs — modern video and world models span diffusion, autoregressive, and hybrid designs, but the center of gravity of each demand function stresses a different part of the silicon stack, and the two should never be summed naively (the hardware consequences are the subject of the companion infrastructure volume). And the **supplier ecosystem** is disjoint: the leaders are Google, ByteDance, Alibaba, Kuaishou, OpenAI (in retreat, as Section 4 documents), and a set of independents, with the enterprise-agent leaders largely absent — which is precisely why token-based telemetry is blind to this economy.

2.3 The integrating frame: two axes of context conquest

The analytical thread that binds generative media to world models and wearables — and binds this report to the rest of the series — is the **Two-Axis Context Map**. Axis 1 is *digital work context*: files, codebases, enterprise systems, conquered through the interface evolution from web chat to IDE to CLI to desktop agent, and monetized against labor budgets. That war — the agentic war — is in its harvest phase, and Reports 3 through 9 of this series measure its economics. Axis 2 is *physical life context*: what people see, hear, and do in the world, conquered through multimodal perception, wearable sensors, and world models, and monetized against attention, commerce, and eventually physical labor via robotics. Axis 2's war is in its sowing phase. Generative media sits at the hinge: commercially, it is a content-production market being repriced today; strategically, the same model families are the entry ramp to world models and embodied AI. The desktop agent is the AI that does your work; the wearable AI is the AI that inhabits your life — high frequency, deep context — and a winner on Axis 2 gains distribution leverage back over Axis 1.

3. The Framework Extended

3.1 Verification at a glance

The enterprise report's foundational law — that the value of generation is gated by the cost and speed of verification — applies with full force here, but the media economy occupies the verification class the enterprise taxonomy touched only in passing: **sensory-parallel human verification**. The human visual system screens an image or video draft in a glance, holistically and in parallel, where a legal memo must be read serially at roughly 200–300 words per minute. The consequence is an iteration economy unlike any text function: visual drafts can be triaged far faster than long-form text, so a creative director can screen a large variant set in the time a research director works through a few pages. Iteration is therefore rapid and cheap on the human side — but a human still sits inside every loop, which caps fully autonomous scaling for most segments.

3.2 The exception that drives the volume: the auction as machine verifier

Two segments partly escape the cap, and the purer of the two is the segment our model makes the volume engine. In performance advertising, the final verifier of a creative asset is not a person but the platform's A/B auction: publish many variants, and click-through and conversion select the winners. The auction is a *scalable statistical verifier* rather than a deterministic one — outcomes are noisy, delayed, and confounded by targeting, auction conditions, and attribution windows — but it can rank large creative sets with far less human review than any other media workflow, playing for advertising a role loosely analogous to the compiler's in software. This is the media economy's structural rhyme with the coding economy, and it predicts the same outcome in structural terms: the machine-verified segment amplifies fastest *per unit of human review*, escaping the bottleneck that caps every other segment. E-commerce SKU video is verified by the same family of commercial metrics and so shares the property, but in weaker form: conversion signals grade a listing's video after the fact rather than selecting among variants as the production logic itself, and the output is templated rather than proliferated. Advertising is therefore the purest and most scalable case of machine verification in media, and it is no accident that ad creative is one of the model's two largest delivered-volume allocations at the opening of the window — running level with the creator economy in 2026 — nor that e-commerce is its companion workhorse. That the creator economy overtakes advertising in delivered minutes from 2027 (Section 5.5) is a breadth-and-depth effect, not a verification effect, and leaves the structural claim intact: creator minutes still pass a human inside every loop, and their growth is bought with human screening time that the auction-verified segment does not spend.

3.3 The monetization-pool cap

The second framework element is the constraint that disciplines the largest volume figures in Section 5: **performance-ad creative and the creator economy draw on the same revenue pool**. A YouTube creator's income is a share of the same advertising and subscription dollars that fund ad campaigns; generated content competes for a fixed pool of audience attention and the ad budget attached to it. Their allocations are therefore *allocative rather than independent* — delivered volume in those two segments is bounded by what the pool can monetize rather than by production capacity, and a further and larger quantity of compute accumulates behind

each delivered minute in exploration and finishing passes that no audience ever sees. The remaining three segments are funded elsewhere — e-commerce from conversion economics, professional production from production budgets, corporate video from operating budgets — so for them the pool operates as a broader economic-consistency screen rather than as a binding ceiling. This is the deepest difference from the enterprise demand function: token demand scales against the labor share of GDP, while the attention-funded core of media demand scales against the attention share of consumer time and spend, a far smaller and slower-growing pool. The cap is also why the model treats capacity and demand as separate questions: the Sora consumer app demonstrated that users will happily consume coding-scale GPU resources when generation is free — one outside estimate placed peak inference expense as high as \$15 million per day against negligible revenue, an estimate built on assumed user activity and per-clip costs rather than disclosed figures, before OpenAI discontinued the consumer app in April 2026 and subsequently announced the shutdown of its API. The shutdown is confirmed; the precise role of inference economics in the decision remains an inference — but the episode frames the category’s constraint well. Capacity to consume compute is proven; what gates *realized* demand is monetization, and our volume assumptions are anchored accordingly to the ROI-measurable segments.

3.4 The unit of account and the shape of the compute

Because diffusion generation is compute-bound and parallel, its natural capacity metric is the GPU-hour, and its unit economics are anchored to throughput: a saturated H100-hour yields roughly three to seven minutes of 1080p video on efficient open-weight models (0.15–0.30 GPU-hours per minute), while premium-tier generation — 4K, physics consistency, phoneme-level lip-sync — runs at 1–5 GPU-hours per minute, so that a single premium 10-second clip costs \$0.50–2.00 in raw compute and consumes the resources of thousands of chat queries. Most delivered media never enters that band: the model applies premium rates only where the consistency requirement demands them, and the intermediate rates that most segments deliver at sit between the two bands rather than inside either. What determines which tier a job runs on is not output resolution — upscaling decouples delivered resolution from generation cost, and every serious model now produces 1080p or native 4K — but the strength of the consistency requirement: the number of denoising steps, the temporal attention range needed to hold a scene together, multi-reference conditioning on characters and products, and native audio. Production routing follows that requirement, and because exploration runs at workhorse quality while only the selected concept is finished at the segment-appropriate delivery-quality rate, the segment-differentiated structure of Section 5 yields a blended average near 0.4 GPU-hours per generated minute. These are the media economy’s equivalents of the tokens-per-task coefficients in the enterprise model, and they carry the same status: explicit, adjustable assumptions disciplined by observed pricing.

4. The Supply Side

4.1 The contenders — and OpenAI's exit

The supply landscape has moved materially since early 2026, and in a single direction: toward more competitors, faster model turnover, and lower prices. OpenAI's retreat from consumer video, provisional a quarter ago, is now all but complete — the Sora web and app experiences were discontinued on April 26, 2026, and OpenAI has announced that the Sora API will follow on September 24, 2026. The withdrawal is consistent with this report's monetization thesis: Sora produced some of the most physically convincing motion in the market yet, by outside estimates, could not cover its inference cost against consumer revenue — one such estimate placed peak spending as high as \$15 million per day, though OpenAI disclosed no figures and never attributed the decision to cost — and the company redirected the capital to the agent war. Video generation is easy to generate and hard to monetize; the company with among the strongest claims to frontier generation quality exited the segment rather than fund it.

What remains is a broadening, genuinely multi-polar field. As of mid-2026 the Artificial Analysis text-to-video leaderboard is led by Google's **Gemini Omni Flash**, which holds the top position in both the with-audio and without-audio categories, while Chinese vendors — ByteDance's Seedance family, Alibaba's HappyHorse, and Kuaishou's Kling — fill most of the remaining top tier. Several highly ranked entries are available only through a single app rather than a broad API — a reminder that a leaderboard score one cannot license is a research result rather than a product.

Gemini Omni Flash and Veo (Google) anchor the Western supply. Gemini Omni Flash, the fast tier of Google's Omni world-model family, currently tops the leaderboard and extends Google's distribution advantage to conversational, quick-turn generation; the Veo line remains the enterprise-safe workhorse — native 4K, strong prompt adherence, reference-based consistency, and 48kHz synchronized *dialogue* rather than merely sound effects — across Fast, Standard, and Quality tiers (\$0.03–0.50 per second) distributed at near-zero acquisition cost through the Gemini app, Google AI subscriptions, and Vertex AI.

Seedance (ByteDance) pairs top-tier quality with price destruction. Seedance 2.0 (February 2026) launched with the market's most permissive input grid — up to nine reference images, three video clips, and three audio files per generation — and the successor 2.5 has extended to native 30-second 4K generation with up to fifty reference images. Its aggressive pricing (on the order of \$0.05 per second on efficient tiers) first made high-volume commercial generation economical; distribution runs primarily through ByteDance's Doubao and Dreamina apps rather than a broad Western API.

HappyHorse (Alibaba) is the consequential new entrant: a model with joint audio-video generation and seven-language lip-sync that ranks second and third on the without-audio board, confirming that the Chinese strength in video is not a two-company artifact but a broad capability grounded in the short-video platforms' data and distribution.

Kling (Kuaishou) competes on cinematic motion and value: native 4K/60fps single-pass generation up to fifteen seconds, multilingual lip-sync, shared audio timelines across multi-shot sequences, at mid-market pricing (~\$0.10–0.14 per second) that undercuts the Western flagships and supplies several of the leaderboard's top-ten entries.

Two structural shifts deserve emphasis. First, the axis of competition has moved off resolution entirely — essentially every serious model now does 1080p or native 4K — and onto clip length, synchronized dialogue, character consistency across shots, and cost. Second, open-weight models have become genuinely useful: Alibaba’s Wan, Lightricks’ LTX-2, and Tencent’s HunyuanVideo now ship 4K-and-audio capability under permissive licenses, the last of these rendering on a single consumer RTX 4090 — which places a floor under the self-hosting economics that inform this report’s GPU-hour assumptions (Section 4.3). Native synchronized audio, a differentiator in 2025, is now a baseline expectation at the model frontier, though its reliability across full production workflows still varies; the frontier has moved to lip-sync precision, shot-to-shot consistency, clip duration, and price.

4.2 Why video is multi-polar where LLMs were not

The LLM market repeatedly resolved toward one or two benchmark leaders; video generation in 2026 is genuinely multi-polar: Google leads the current quality rankings, but ByteDance, Alibaba, and Kuaishou occupy most of the remaining top tier — a depth of Chinese co-leadership with no LLM-market parallel. The structural reason is the vertical integration of data and distribution: ByteDance and Kuaishou are the world’s largest short-video companies, possessing both unusually large proprietary short-video datasets and distribution feedback (billions of hours of the exact content type being generated) and the channels that monetize the output. In multimodal generation, China is not a fast follower but a co-leader — a fact with direct consequences for the cost curve, since Chinese price competition is the principal force collapsing unit costs, and for measurement, since Seedance distributes primarily through ByteDance’s Doubao ecosystem without a public Western API, limiting independent verification of its usage claims. Market behavior mirrors the agent market’s multi-homing: professionals pick the model per shot through aggregator platforms, and leaderboard rankings reshuffle on a monthly cadence.

4.3 The cost collapse, and the status of the GPU-hour coefficients

Retail generation pricing now spans roughly \$0.03–0.50 per second across tiers — a workhorse band near \$0.05–0.15 (Seedance efficient tiers, Kling value tiers, Veo Fast, Gemini Omni Flash) and a premium band above it — while self-hosted open-weight inference runs an order of magnitude lower. At current workhorse rates a ten-minute first-pass generation costs roughly \$30–90 before retakes and editing; under the further five- to ten-fold price decline this report models, that falls to approximately \$3–18 by 2028 — below a stock-footage subscription, the threshold past which generation becomes the default rather than the option (the cost curve of Section 5.8). The first-order losers are stock-footage libraries, small production houses, and the lower tiers of the advertising production chain; the first-order winners are the platforms that monetize distribution and the suppliers of inference compute — video generation being orders of magnitude more compute-intensive per unit of output than text.

A note on the status of the GPU-hour coefficients that convert generated minutes into compute demand in Section 5. The workhorse range of 0.15–0.30 GPU-hours per generated minute — equivalent to roughly three to seven minutes of video per saturated accelerator-hour — and the delivery coefficients of 0.30 to 3.00 for finished work — rising with the consistency requirement, and reaching the premium band only for a creator’s hero shot (1.20) and for professional 4K with physics consistency and phoneme-level lip-sync (3.00) — are **explicit**

engineering assumptions, informed by published model throughput and by the retail and managed-API price points above, but not directly identifiable from them: service prices bundle hardware choice, batching, utilization, subsidy, and provider margin, none of which vendors disclose. Managed-API pricing at cents per clip is *consistent with* low underlying generation costs; it cannot pin the coefficient. The resulting GPU-hour totals should therefore be read as conditional capacity scenarios under stated intensity assumptions rather than as measured demand, with the residual risk skewed in one direction — continued efficiency gains and open-weight maturation would bias the coefficients downward, lowering the GPU-hour totals without changing the generated-minute volumes. Because retail prices fall faster than underlying compute cost per minute (margin compression plus efficiency), the demand-side cost collapse driving volume and the supply-side GPU-hour intensity are partly decoupled — which is why falling prices expand generated minutes without proportionally shrinking the fleet the segment requires.

5. The Functional Anatomy of Media Demand

This is the anatomical core of the report. The media track is not one demand curve but five, each with a distinct function, a distinct content requirement, a distinct customer, a distinct platform owner, and — decisively — a distinct demand-shape: some segments multiply generated **minutes** while others barely register in minutes yet dominate **GPU-hours**. This chapter builds each segment forward from a stated reference anchor in the existing online-video economy, states the penetration logic explicitly, and then screens the resulting total against the monetization constraint of Section 3.3, which operates as a qualitative economic consistency check on the volume assumptions rather than as a computed ceiling.

5.1 Method: anchor, penetrate, deepen — then screen against the pool

Four of the five video segments are estimated in two stages. The first stage produces **delivered minutes** — the minutes that actually reach an audience — as *installed base × AI-adoption penetration × output per adopter*. The second stage converts delivered minutes into compute through an explicit two-pass intensity:

$$\text{GPU-hours per delivered minute} = \text{exploration passes} \times \text{workhorse rate} + \text{finishing passes} \times \text{delivery rate}$$

This is the structure practitioners work in, and the separation is not cosmetic. Creators and studios do not render every attempt at final quality. They explore widely and cheaply — storyboard frames, low-resolution motion tests, alternative takes of a single shot — and then render the selected concept once or twice at whatever quality the finished piece requires. Pricing every generated minute at a single blended rate therefore charges delivery-quality rates for drafts that were never rendered at delivery quality, and it does so most heavily in the segments that explore hardest.

Two properties of the two-pass structure govern every segment below. **Exploration depth per task is held flat across the forecast** rather than escalating year by year. What bounds it is not

generation cost, which is collapsing, but the human review bandwidth needed to look at the results: a creative director who can screen fifty candidates today will not screen five hundred in 2028 because each one became ten times cheaper. This is the media analogue of the review-bandwidth cap that governs qualitative work in the enterprise model. The flat assumption applies to depth per task, not to a project's total exploration volume: where AI spreads into a larger number of shots and workflow stages — the professional case — total exploration per project rises even though the number of candidates reviewed for any one shot does not. **Finishing passes are held near-constant at 1.5 to 2** across every segment, independent of how deep the exploration runs. Only the selected concept is rendered for delivery, and in practice it rarely emerges exactly to plan on the first attempt — a small defect in a hand, a lip-sync drift, a lighting mismatch — so one to two finished renders per delivered minute is the working norm regardless of whether exploration ran to four passes or twenty-five. The *rate* at which those renders run is a separate question, and it varies by an order of magnitude across the five segments: it is set by the consistency requirement of Section 3.4, not by a universal premium tier. A templated e-commerce clip is delivered barely above workhorse cost (0.30), a performance ad or a corporate module only modestly above it (0.40), a creator's hero shot at mid-tier (1.20), and professional footage at the full premium rate (3.00). Only the last two reach the premium band described in Section 3.4; the other three are finished at rates that sit close to exploration cost, which is why those segments' effective intensities stay near or below one GPU-hour per delivered minute.

The exploration depths below are calibrated against published production accounts rather than vendor claims, because vendors have an evident interest in presenting their models as first-time-right: documented workflows report roughly two to four generations per finished clip for light social work, around four for narrative work with character consistency, and substantially higher ratios for ad-grade and broadcast-grade material, where ten to forty prompts per final clip is commonly described. Where a segment's evidence spans a range, the upper end is taken. These ratios are clip-level: they describe how many candidates are generated for one finished clip. They inform the exploration assumptions in the four delivered-minute segments, where each segment's own verification logic then adjusts them — advertising sits below the ad-grade range because the auction, not a director, does the selecting, and variants are published rather than discarded. The professional segment adopts the upper end of the range directly, but at the level of the individual shot rather than the delivered minute, through the derivation in Section 5.6.

Professional media and entertainment is the exception, and Section 5.6 sets it out in full. Because the great majority of that segment's compute is consumed by project-level exploratory work — concept development and pre-visualization that has no one-to-one relationship with anything that reaches the screen — its exploration workload is estimated directly per adopting project, and its delivered footage is modeled separately alongside it. The per-delivered-minute figure quoted for that segment in Table 1 is therefore a derived diagnostic, not an input.

The denominators are reference anchors and scenario constructions of differing evidentiary strength rather than uniformly observable statistics: Meta's officially disclosed base of more than 10 million advertising businesses, widened to ~15 million *platform-advertiser accounts* across the major ad platforms (an account basis that tolerates cross-platform duplication, since campaigns generate creative per platform); the ~2 billion live product listings of the

global marketplaces; YouTube’s upload firehose (~500 hours per minute, roughly 15.8 billion published minutes a year) and its creator hierarchy (~3 million monetizing channels inside a ~15 million weekly-cadence creator population); an order-of-magnitude construction of ~60,000 professional productions completed annually (features, scripted-television episodes, music videos, premium documentaries); and a scenario denominator of ~4 million organizations above fifty employees. The volume coefficients carry wide error bands, and the sections that follow state the confidence attached to each; every input is explicit and adjustable.

Table 1. Delivered video volume by segment (millions of minutes reaching an audience)

Video segment	Base (anchor)	Penetration 26→28E	2026E	2027E	2028E
1. Performance ad creative	~15M platform-advertiser accounts	15% → 55%	54	144	330
2. E-commerce SKU video	~2,000M live marketplace listings	5% → 28% of stock (~35% → ~88% GMV-weighted)	43	138	280
3. Creator economy / UGC	~15M weekly-cadence creators	12% → 40%	56	234	624
4. Professional M/E	~60K professional productions/yr	25% → 75% (any-use)	0.15	0.54	1.1
5. Corporate internal	~4M organizations (50+ employees)	5% → 30%	16	54	120
Video total delivered			169	571	1,355

Table 1b. Compute-intensity structure, and generated video volume by segment (millions of minutes actually produced)

Video segment	Exploration workload × workhorse rate	Finishing passes × de- livery rate	GPU-hr per delivered minute	2026E	2027E	2028E
1. Performance ad creative	4.0×0.25	1.5×0.40	1.60	297	792	1,815
2. E-commerce SKU video	2.0×0.15	1.5×0.30	0.75	151	483	980
3. Creator economy / UGC	5.0×0.25	2.0×1.20	3.65	393	1,638	4,368
4. Professional M/E	project-level: ~1,700 → ~5,550	2.0×3.00	57 → 73 (derived)	26	121	252

Video segment	Exploration workload × workhorse rate	Finishing passes × de- livery rate	GPU-hr per delivered minute	2026E	2027E	2028E
	exploration min/ title × 0.30					
5. Corporate internal	2.0 × 0.20	1.5 × 0.40	1.00	56	189	420
Video total generated				922	3,223	7,835

Generated minutes are delivered minutes multiplied by the total pass count — exploration plus finishing — and are the volume the models actually produce; delivered minutes are the subset that reaches an audience. Professional M/E is the exception noted in Section 5.6: its exploration is set per project rather than per delivered minute, so its per-minute intensity is derived from the two figures rather than assumed.

Evidentiary status of the inputs. The bases divide into three grades, and readers should weight them accordingly. **Observable anchors** (grade A–B): Meta’s disclosed advertiser base, YouTube’s published upload rate and monetizing-channel count, Microsoft and Google seat and customer disclosures. **Constructed anchors** (grade C): the ~15 million platform-advertiser accounts widened across platforms, the ~15 million weekly-cadence creator population, the ~2 billion marketplace listings, the ~60,000 annual professional productions, the ~4 million organizations above fifty employees — each assembled from public components rather than reported as such. **Scenario assumptions** (grade D): every penetration path, every output-per-adopter figure, and the exploration and finishing coefficients, including the professional segment’s shot counts and candidate ratios. Nothing in the last group is a measurement, and the totals inherit that status.

Three movements across the two tables deserve immediate notice. Delivered minutes expand eight-fold over two years, and generated minutes expand with them, because breadth and depth compound. The volume leadership is contested at the start and settled by the end: advertising and the creator economy are within a few percent of one another in 2026 delivered minutes, but the creator economy pulls decisively ahead from 2027 and holds 46% of delivered minutes by 2028 — advertising’s growth is bounded by advertiser count and fatigue-cycle economics, while creator demand compounds on both breadth and depth. And the intensity column of Table 1b, not the minutes column of Table 1, is where the compute ranking is decided: professional media and entertainment delivers a rounding error in minutes but carries more than seventy GPU-hours behind each one, while e-commerce delivers two hundred and fifty times as many minutes at roughly one-hundredth the intensity. The professional row is also the one segment whose intensity is derived rather than assumed: its exploration is set at project level and divided through by delivered minutes, for the reasons given in Section 5.6. A third reading belongs to Table 1b alone — the ratio of generated to delivered minutes, which runs from 3.5 in e-commerce to 5.5 in advertising, 7 in the creator economy, and above 200 in professional production. That ratio is the exploration structure of Section 5.1 made visible,

and it is why compute rankings and volume rankings do not agree. The sections below derive each row.

5.2 Why the penetration rates differ: denominators, deciders, and forcing functions

Read naïvely, the penetration column is counter-intuitive: professional studios — the most quality-sensitive, rights-sensitive customers in the market — carry the highest rate, while e-commerce and corporate video, seemingly the easiest use cases, carry the lowest. The paradox dissolves once three structural variables are made explicit, because the five rates are defined over different kinds of denominator and driven by different adoption mechanics.

The first variable is the **nature of the denominator** — whether it is a *flow* of live projects that all make a tooling decision this year, or a *stock* that includes a dormant long tail which no rational actor would ever address. The second is the **decision structure** — whether adoption is decided by a few thousand concentrated professional organizations or by millions of heterogeneous small businesses and individuals. The third is the presence or absence of a **forcing function** — whether the platform injects the technology from the supply side, or competitive cost pressure compels it, or adoption is a purely voluntary pull that requires creating a workflow that did not previously exist.

Segment	Denominator	Decision structure	Forcing function	Resulting pace
4. Professional M/E	Flow — every unit is a live, budgeted project	Highly concentrated (thousands of studios/VFX houses)	Competitive cost pressure (five-figure VFX shots vs. cents)	Fastest
1. Ad creative	Near-flow (active advertisers)	Dispersed, but the platform decides on their behalf	Supply-side injection (Advantage+ / PMax auto-generation)	Fast
3. Creator economy	Flow (weekly uploaders)	Fully dispersed individuals	None — cost and quality pull only	Moderate
5. Corporate internal	Stock — most organizations have no video function today	Dispersed, plus procurement friction	Weak — suite bundling lowers the friction but a new workflow must still be created	Slow
2. E-commerce SKU	Stock — dominated by a dormant listing tail	Platform-concentrated, but application is traffic-weighted	Partial (platform auto-conversion)	Lowest on stock; highest on GMV

The rest of this chapter applies this framework row by row; the apparent paradoxes in the table are resolved cases of it.

5.3 Segment 1 — Performance Ad Creative: supplier-push, not customer-pull

Function, content, customer. Short performance-marketing video — feed ads, Reels, Shorts, and TikTok placements — typically eight to fifteen seconds, produced in large variant sets for continuous A/B testing. The requirement is volume and freshness rather than cinematic polish, so exploration runs entirely at the workhorse rate (0.25 GPU-hours per minute) and finishing is only lightly upgraded (0.40). The customer base is the ~15 million platform-advertiser accounts across the major platforms (anchored on Meta’s officially disclosed base of more than 10 million advertising businesses, plus the Google, TikTok, and Amazon account bases, counted with cross-platform duplication since creative is generated per platform), overwhelmingly SMBs for whom video production was previously uneconomic; a \$10 generation budget now yields dozens of usable clips, dropping video creation to roughly the cost structure of copywriting.

The verification logic. This is the purest case in media of a workload whose final verifier is machine-mediated rather than human — the platform’s A/B auction: publish many variants and let click-through and conversion select winners, with human review shrinking to brand-safety screening. As Section 3.2 argued, the market plays the role the compiler plays in software — and creative-fatigue cycles (ad sets decay in two to four weeks on short-form platforms) then make continuous refresh rational, so that when the marginal cost of a variant approaches zero the profit-maximizing variant count rises.

The penetration logic (15% → 55%). The steepest non-professional ramp in the model rests on the observation that adoption here is not a customer decision but a platform decision. Meta disclosed as early as September 2024 that more than one million advertisers were using its generative ad tools, generating fifteen million ads in a single month; Meta’s Advantage+ and Google’s Performance Max increasingly generate video variants server-side from an advertiser’s existing assets, which means that by 2028 the operative definition of penetration shifts from “adopted” to “did not opt out.” The 55% ceiling reflects the residual: categories adequately served by static imagery, regulated verticals (pharmaceutical, financial) whose creative-approval chains block automated variants, and the low-activity advertiser tail that supplies no usable source assets. Our 2026 figure of 2.25 million *video-specific* AI advertisers is a deliberate subset of the larger disclosed any-generative-AI base.

Parameters. Advertisers using AI video rise 2.25M → 4.5M → 8.25M; published AI clips per advertiser rise 120 → 160 → 200 per year on the fatigue-refresh logic; average published clip ~12 seconds (0.2 minutes). Product: **54 → 144 → 330 million delivered minutes**. Intensity is four exploration passes at the workhorse rate plus 1.5 finishing passes at 0.40 — a lower exploration depth than the creator and professional segments, because the auction, not a creative director, does the selecting: variants are the deliverable rather than candidates for discard, so relatively little of what is generated is thrown away. At 1.60 GPU-hours per delivered minute the segment consumes **86 → 230 → 528 million GPU-hours**.

One layer above this segment is not counted. Agencies and brand marketing teams use video models heavily, but for a different purpose: idea exploration. A single campaign concept may be sketched as hundreds of short clips, shown internally and to the client, and discarded — with the delivered commercial then shot conventionally and finished with AI-assisted post-production. Compute is consumed at volume while almost nothing generated reaches air,

and the exploration depth involved plausibly exceeds even the professional-production figure in Table 1. Anchoring that demand would require a campaign-level base population that no published source supports, so it sits outside this estimate — and its absence biases the advertising figures downward by a margin that is not small.

The platform reading. Meta and Google own this segment three times over: the auction that verifies the creative, the advertising pool that monetizes it, and increasingly the generation model that produces it. The supply-side injection that drives the penetration curve exists precisely because of that vertical integration — a platform that owns both the model and the monetization pool has every incentive to generate the creative itself, a dynamic this series' analysis of consumer agentic platform competition among Google, Meta, and Amazon develops at length.

5.4 Segment 2 — E-commerce SKU Video: lowest penetration of stock, highest penetration of commerce

Function, content, customer. Conversion of static product listings into short demonstration and lifestyle clips — five to ten seconds, templated, lightly retaken — verified by the listing's own conversion rate, hence machine-graded and run at the cheapest tier in the model (0.15 GPU-hours per minute for exploration, 0.30 for the finished clip). The decision-relevant customers are the platforms and their seller bases: Amazon's ~1.6 million active sellers and its advertising console (whose AI video generator has been offered to advertisers at no charge since 2024), Shopify's ~4.9 million stores, and the Chinese marketplaces, several of which already run listing-image-to-video conversion at platform scale. The *unit*, however, is the SKU: roughly 2 billion live listings globally (the weakest-grade anchor in the model), of which Amazon alone accounts for several hundred million.

The penetration logic (5% → 28% of stock; ~35% → ~88% of GMV). The intuition that e-commerce should be the most-penetrated segment is correct — the low figure in the table is a denominator effect, not a slow-adoption claim. Video creates value only where there is traffic to convert, so the economically rational application order is GMV-weighted from the top: under the marketplace-typical Pareto structure in which the top decile of listings carries roughly 85% of merchandise value, our 2028 stock coverage of 28% decomposes as the entire top decile plus roughly a fifth of the dormant tail — which is approximately **88% of GMV-weighted commerce** ($85\% + 0.2 \times 15\%$). On the measure that matters commercially, this segment ends the forecast as the *most* penetrated function in the taxonomy; the stock-based figure is retained in the table because GPU-hours are generated per listing, not per dollar. One reconciling note: the 2026 GMV figure of ~35% embeds imperfectly traffic-ordered early application — platform auto-conversion sweeps mid-tail listings alongside the head — rather than a strict top-down fill, which at 5% stock coverage would imply roughly 42% of GMV. The dormant tail below the traffic line is not late-adopting; it is rationally never addressed.

Parameters. Covered SKUs rise ~100M → ~300M → ~560M (5% → 15% → 28% of stock); annual delivered minutes per covered SKU run 0.43 → 0.46 → 0.50 (two to three published clips of ~10 seconds). Product: **43 → 138 → 280 million delivered minutes**. Intensity is the lowest in the model — two exploration passes at 0.15 plus 1.5 finishing passes at 0.30, for 0.75 GPU-hours per delivered minute — because the output is templated, the clips are short, and the verifier is a conversion rate rather than a human eye, so there is little to explore and little to perfect.

The segment consumes **32 → 104 → 210 million GPU-hours**. Amazon is this segment's platform owner in the West on the same triple logic as Meta in advertising — it owns the traffic, the conversion data that verifies the creative, and now the generation tooling — with Taobao and Douyin commerce playing the equivalent role in China.

5.5 Segment 3 — Creator Economy / UGC: the depth-driven leader from 2027

Function, content, customer. The individual and small-business creator economy: YouTube channels, short-form operators, the video-podcast wave that Spotify and YouTube are both pushing, and independent musicians. Content spans B-roll and intros today, graduating toward assembled segments and AI-hosted formats; exploration runs at the workhorse rate of 0.25 GPU-hours per minute and finishing at a mid-tier 1.20, reflecting the native audio and character consistency that hero shots increasingly require. The base is the ~15 million creators publishing on a weekly cadence — anchored on YouTube's ~3 million monetizing channels and widened to regular non-monetizing uploaders, video podcasters, and short-form businesses, while deliberately excluding the ~65 million casual-creator figure whose long tail publishes too irregularly to matter for volume.

The penetration logic (12% → 40%): no forcing function, and a genre-fit ceiling. Adoption here is fully dispersed individual decision-making with no platform injection, so the curve is pure cost-and-quality pull — and its ceiling is set by genre structure rather than by willingness. A large share of creator output is in genres where AI generation is structurally inapplicable or only marginally applicable: gaming (the screen capture is the content), vlogging and IRL formats (authenticity is the product), reaction and commentary. Treating the high-applicability genres — education and explainers, ambient and music channels, kids' content, listicle and documentary-style formats — as roughly 50–60% of serious-creator output, and assuming adoption within that applicable pool reaches ~70% by 2028, yields the 40% aggregate: the headline rate is the product of genre fit and within-genre adoption, not a free-hand S-curve. The 2026 figure of 12% reflects that the tooling is already trivially accessible to this population and that platform integration is underway — YouTube has embedded Veo-powered generation directly into Shorts — but that full-video workflows remain immature.

Why content depth, not creator count, drives growth. The creator base is roughly stable and upload cadence is fixed by format (~52 videos per creator-year), so this segment's growth is carried by adoption and by content depth rather than by head-count. Delivered AI minutes per video rise 0.6 → 1.2 → 2.0 as the technology traverses the insert-to-segment-to-full-generation progression of Section 5.8. The product — $15\text{M} \times 12\%/25\%/40\% \times 52 \times 0.6/1.2/2.0$ — yields **56 → 234 → 624 million delivered minutes**, which overtakes advertising in 2027 and reaches 46% of all delivered minutes by 2028. Exploration depth is not the second growth term but the amplifier applied to both: five workhorse passes per delivered minute in every forecast year, a depth consistent with documented narrative workflows where roughly one generated clip in four survives the edit and the surviving clip is itself trimmed. It sets the level of this segment's compute rather than its slope. Two finishing passes at 1.20 reflect the mid-tier rendering that hero shots increasingly receive as native audio and character consistency become table stakes. At 3.65 GPU-hours per delivered minute this segment consumes **205 → 854 → 2,278 million GPU-hours**, roughly 70% of the video track by 2028 — the single largest concentration of compute in the media demand function. A scale check disciplines the optimism: 624 million delivered AI minutes in 2028 is roughly 4% of the current annual upload volume of YouTube

and the short-form platforms combined, comfortably inside any plausible slop-saturation ceiling.

The platform reading. Google is this segment's vertically integrated owner: YouTube supplies the distribution, the Partner Program supplies the monetization pool (a share of the same advertising dollars that fund Segment 1 — the double-counting the cap of Section 3.3 exists to prevent), and Veo, embedded in the creation surface, increasingly supplies the generation itself. No other player holds all three layers of any segment this completely.

5.6 Segment 4 — Professional M/E: the highest headline penetration and the Netflix counter-intuition

Function, content, customer. Studios, streamers, and music-video production: VFX shots, pre-visualization, concept work, B-roll, and increasingly generated final footage, at the top of the quality ladder — 4K mastering, physical consistency, phoneme-level lip-sync — with finishing modeled at 3.0 GPU-hours per minute, ten times the workhorse rate at which this segment's very deep exploration runs. The customer set is tiny: roughly 60,000 professional productions a year globally (on the order of 9,000 feature films, 35–40,000 scripted-television episodes, and the music-video and premium-documentary flow).

The penetration logic (25% → 75%, any-use): why the most demanding customer adopts fastest. The rate that most offends intuition is the most structurally over-determined in the model, for three reasons drawn directly from the Section 5.2 framework. First, the denominator is a pure *flow*: every one of the 60,000 units is a live, budgeted project whose team makes toolchain decisions this year — there is no dormant stock diluting the base, unlike the listings tail or the video-less corporate majority. Second, the *any-use threshold* is low and nearly risk-free: a production counts as penetrated if AI generation appears anywhere in its pipeline, and the earliest entry points — pre-visualization, concept frames, temp B-roll — never reach the final screen, so the hallucination-and-rights risk that gates final footage does not gate adoption itself, while the cost differential is immediate and large — with the caveat that this compares finished-shot budgets against raw generation compute, before labor, compositing, quality control, and rights costs, so it is directional rather than like-for-like. Third, the decision structure is concentrated: a few thousand studios, VFX houses, and production companies under identical competitive cost pressure diffuse innovations at professional-market speed, the same dynamic that made AI coding tools near-universal among professional developers within roughly two years (84% using or planning to use them, per Stack Overflow's 2025 survey). The public record confirms the mechanism is already operating: Lionsgate's partnership with Runway and Netflix's disclosure of its first generative-AI final footage in a released title mark the majors' transition from experimentation to production use. The economic substance of this segment, however, is carried not by the penetration rate but by exploration volume per adopting title.

Why this segment is anchored on projects rather than on screen time. Studio use of generative video does not resemble the other four segments, and modeling it from screen time alone would misstate it in both directions. What a production consumes is overwhelmingly *exploratory*: concept frames, pre-visualization passes, motion tests, alternative takes of a single shot — generated at low resolution, evaluated in a review session, and discarded. This is idea communication, not footage, and its volume is set by the project's creative uncertainty, not by

the length of the finished piece. What eventually reaches the screen is a much smaller quantity of AI-touched material, and it is typically AI-assisted post-production on conventionally shot footage — a replaced background, a de-aged face, an extended set — rather than wholly generated scenes.

The segment is therefore built from the project rather than from the minute. Exploration volume per adopting title is derived at the shot level: roughly forty candidate clips of around eight seconds for each AI-assisted shot or sequence — the upper end of the ten-to-forty range that published production accounts report for broadcast-grade work, applied here because studio review sessions are exactly where large candidate sets are screened cheaply — across a number of AI-assisted shots and sequences that rises from ~300 to ~700 to ~1,000 per title as AI use widens from pre-visualization into VFX development and scene work. That is ~1,700 → ~4,000 → ~5,550 exploration minutes per title, generated at the workhorse rate. Note what does and does not grow: the candidate count per shot is flat, consistent with the review-bandwidth argument of Section 5.1; what rises is the number of shots AI is asked to explore. To this is added ~10 → ~18 → ~25 delivered AI-touched minutes per title, rendered twice at the premium rate of 3.00. That yields roughly 570 → 1,310 → 1,815 GPU-hours per adopting title, and **9 → 39 → 82 million GPU-hours** for the segment. The delivered figures are deliberately bounded by runtime: the base averages well under an hour per unit once 45-minute episodes and music videos are counted alongside features, so 25 minutes of AI-touched footage reaching the screen in an adopting title is already a generous reading. Because exploration accounts for roughly nine-tenths of the segment's compute, the total is insensitive to the delivered-footage input and highly sensitive to the count of AI-assisted shots per title — which is the number a reader wishing to stress this segment should vary.

The counter-intuition, made precise. The instinct that streaming and film drive media-AI demand is false on a minutes basis and only modestly true on a compute basis. The segment delivers **0.15 → 0.54 → 1.1 million AI-touched minutes — under one tenth of one percent of the track** — and consumes **9 → 39 → 82 million GPU-hours, roughly 2.5% of video GPU-hours**. It is thus over-represented in compute by a factor of about thirty relative to its delivered volume, which is the durable form of the counter-intuition: premium intensity and very deep exploration make this by far the most compute-hungry delivered minute in the media economy, without making it the compute-dominant segment. Localization — one title regenerated across a dozen or more dubbing languages — adds a further multiplier and is the principal B2B driver of the voice track in Section 5.10. The broader lesson survives intact: the segment nearly invisible in volume statistics punches far above its weight in the compute bill, and analysts reading delivered minutes as a compute proxy will misread the market exactly as naïve token counts misread the LLM market.

A note on what this segment does not include. The ~60,000 base counts scripted premium production — features, scripted television episodes, music videos, premium documentaries. It does not count day-to-day broadcast operations: the news, sports graphics, weather visualization, and educational programming that thousands of broadcasters produce daily, and in which AI-generated explanatory footage is a visibly growing practice. That population is larger than the scripted flow by orders of magnitude, but its per-unit intensity is far lower — short lead times cap exploration at two or three passes, and the quality bar is explanatory rather

than cinematic — so an illustrative sizing places it below one percent of video GPU-hours. It is excluded for want of an anchorable base rather than for want of significance.

5.7 Segment 5 — Corporate Internal: the structural laggard, honestly sized

Function, content, customer. Internal learning-and-development modules, corporate communications, policy explainers, and presentation video, replacing template stock footage and talking-head shoots. Avatar-led synthesis dominates the workload, and because avatar pipelines are computationally lighter than full-scene diffusion, the segment runs at 0.20 GPU-hours per minute for exploration and 0.40 for the finished module. The base is the ~4 million organizations above fifty employees.

The penetration logic (5% → 30%): why the “easiest” use case moves slowest. This is the mirror image of Segment 4 on all three framework variables. The denominator is a *stock* in which most organizations have no video-production function today — no budget line to redirect, no incumbent cost to undercut. The forcing function is weak rather than absent: suite bundling lowers procurement friction, as the rest of this paragraph describes, but a company that produces no internal video still loses nothing visible by continuing not to, and adoption requires creating a workflow rather than substituting within one. And the decision path runs through procurement and IT approval, the highest-friction channel in the taxonomy. The 2026 starting point is disciplined against vendor reality: the category leader’s disclosed customer base sits near 50,000 enterprises (including most of the Fortune 100), and the dedicated-vendor installed base across the category plausibly totals 100–200 thousand organizations. Our figure of 200,000 adopting organizations (5%) therefore embeds one further assumption, which recent product history supports: video generation is ceasing to be a procured product and becoming a bundled feature of software organizations already own — Google Workspace now ships Vids with Gemini generation, Canva has folded AI video into the most widely deployed design tool in the mid-market, and Microsoft 365’s trajectory points the same way. Bundling removes precisely the procurement friction that makes this segment slow, which is what carries the curve from a vendor-anchored 5% to 30% by 2028 — still the second-lowest endpoint in the table, because bundled availability creates capability, not the communications practice that uses it.

Parameters. Adopting organizations rise 0.2M → 0.6M → 1.2M at ~80 → ~90 → ~100 delivered minutes per organization-year: **16 → 54 → 120 million delivered minutes**. Intensity is 1.00 GPU-hour per delivered minute — two exploration passes at 0.20 plus 1.5 finishing passes at 0.40 — the avatar-led pipeline being both computationally lighter than full-scene diffusion and, because the script is fixed and the presenter synthetic, unusually low in exploration. The segment consumes **16 → 54 → 120 million GPU-hours**. This is the segment furthest from the attention pool, funded from operating budgets rather than from advertising or subscription revenue, and it is modeled conservatively for exactly that reason: its constraint is organizational habit, not monetization.

5.8 The three compounding curves behind every path

Each segment’s year-over-year trajectory is the product of three curves stated here once. (a) *Technical completeness — insert → segment → full generation*. The typical 2026 production envelope — still dominated by reliable 5–15-second clips even as frontier models reach roughly thirty seconds, with face-and-hand consistency unstable, multi-shot windows around fifteen

seconds, and native audio standard at the model frontier but uneven across production-grade workflows — confines current practice to inserting AI B-roll into filmed video; as character consistency, minute-scale scenes, and audio synchronization mature, 2027 supports generating roughly half of a piece and 2028 supports fully generated formats. For the four delivered-minute segments this primarily raises delivered minutes per piece rather than exploration depth, and it is why the creator segment — content-depth-driven rather than count-driven — accelerates hardest; in professional M/E the same maturation also widens the number of shots and workflow stages subjected to exploration, which is where that segment’s growth comes from. (b) *Cost — crossing the stock-footage line.* Retail generation spans roughly \$0.03–0.50 per second (workhorse tiers near \$0.05–0.15), with self-hosted open-weight inference at \$0.007–0.022; a ten-minute first-pass generation therefore costs roughly \$30–90 today, and the modeled further 5–10x decline takes it to approximately \$3–18 by 2028 — at the lower end, below a stock-footage subscription, past which generation becomes the default rather than the option. That comparison is a first-pass figure: a delivered minute in a segment that explores deeply costs several times the headline, which is why the cost line is crossed in different years by different segments. (c) *Social acceptance — from “AI slop” to normalization.* The 2025–26 viewer backlash and platform labeling mandates are treated as transitional friction that relaxes as quality crosses detection thresholds and music-licensing disputes settle into royalty structures; acceptance operates on the *ceilings* of the penetration curves above, and it is the model’s most falsifiable assumption (Sections 7–8).

5.9 The platform synthesis: three pool-owners, one two-axis competitor

Mapped against the platform war among Google, Meta, and Amazon that this series analyzes on the consumer side, the five segments resolve into spheres of vertical integration. Meta and Google own performance advertising (Segments 1 and, for Google, the search-and-YouTube ad surfaces) — the auction, the ad-dollar pool, and now the server-side generation. Google separately owns the creator economy through YouTube — distribution, the Partner Program’s monetization pool, and Veo embedded in the creation surface. Amazon owns Western commerce video — the traffic, the conversion telemetry, and free generation tooling for its sellers and advertisers — with the Chinese platforms playing the equivalent integrated role in their market. The pattern is uniform: **supply-side injection, the strongest forcing function in the penetration framework of Section 5.2, appears exactly where a platform owns both the generation model and the monetization pool.** It is also asymmetric in one direction that the rest of this report keeps encountering: Google alone appears in three of the five segments (advertising, creator, and — through Workspace Vids — corporate internal), before counting the world-model and wearable positions of Section 6. The media demand function, read at the platform layer, repeats the same structural pattern: no other single company currently spans as many of these segments within one ecosystem.

5.10 The image and audio tracks

Two non-video tracks complete the demand function, dominated by the same customers. **Image generation** — advertising visuals and e-commerce product imagery leading commercial volume — scales ~50 → ~150 billion units at ~0.001 GPU-hours each, contributing 50 → 150 million GPU-hours. **Voice and music** — TTS narration for the creator and corporate segments, AI background music, and the dubbing-localization demand that Segment 4 generates — scales ~0.5 → ~2.5 billion output hours at ~0.01 GPU-hours per hour, contributing 5 → 25 million. Both

are workhorse-priced and conversion- or sensory-verified; neither reshapes a total that video dominates throughout.

5.11 Totals and the fleet equivalent

Table 2. Media-track GPU-hours and fleet equivalent (conditional mid-case scenario under the stated intensity assumptions)

Item	2026E	2027E	2028E
Video: performance ad creative (M GPU-hours)	86	230	528
Video: e-commerce SKU (M GPU-hours)	32	104	210
Video: creator economy / UGC (M GPU-hours)	205	854	2,278
Video: professional M/E (M GPU-hours)	9	39	82
Video: corporate internal (M GPU-hours)	16	54	120
Video subtotal (M GPU-hours)	348	1,281	3,217
Image generation (M GPU-hours)	50	90	150
Voice & music (M GPU-hours)	5	12	25
Total (M GPU-hours)	403	1,383	3,392
Fleet equivalent ('000 H100-equivalent accelerators, at segment-specific utilization)	108	379	944

A word on the utilization rates behind the fleet row, which differ by segment rather than applying one figure to the track. Diffusion generation is compute-bound batch work with no latency SLA, so enterprise pipelines that queue jobs — e-commerce conversion, corporate video, studio rendering — sustain roughly 55% utilization, materially higher than the interactive agent loops of the LLM model. Self-serve advertising runs near 45%. Individual creators are the exception at 38%: they generate inside interactive sessions, waiting for each result before prompting again, which is a latency-bound pattern much closer to customer-facing conversational workloads than to batch rendering. Because the creator segment carries most of the track's compute, that single rate does more to set the fleet total than any other utilization assumption.

The trajectory is an eight-fold expansion in GPU-hours over two years — against the roughly two-fold growth of compute-weighted B2B token demand in the companion enterprise model. The fleet figures are H100-equivalent units; production systems mix GPUs, TPUs, and ASICs, so the row measures compute scale, not a chip count. Two cautions attach. This is a conditional mid-case scenario, not a point forecast — plausibly varying by several-fold given the limited observability of production throughput and creator adoption — screened against the advertising-and-subscription monetization pool of Section 3.3 from above and anchored by the segments already in production from below, with the dominant sensitivities being creator adoption (12 → 40%), creator delivered minutes per video (0.6 → 2.0), and the exploration depth assumed for the creator segment, which together set roughly 70% of the total; the professional segment's exploration volume per title is the next most consequential input, though it governs a far smaller share. And the two demand functions should be treated as independent by investors: they respond to different price curves (visual-generation unit costs collapsing on

Chinese competition), different supplier sets, and different chip preferences — sustained dense compute for the diffusion-centered media stack against memory bandwidth for autoregressive decode, subject to the center-of-gravity caveat of Section 2.2 (several leading video and world models are themselves autoregressive or hybrid) — a divergence whose hardware consequences the infrastructure volume develops.

6. Beyond Content

6.1 From video models to world models: Google’s integrated path

Video models predict plausible next frames; world models predict how an environment evolves *in response to actions*. The two sit on one technical continuum — a world model is video generation plus physical consistency plus interactivity — and Google currently operates the most visibly integrated position along its full length — the Omni family’s fast tier anchoring the consumer video market of Section 4 at one end, the dedicated Genie line examined here at the other — though not alone on any single stretch of it: NVIDIA’s Cosmos family positions open foundation models for physical AI, and Meta, World Labs, Waymo, and a range of robotics developers compete on individual layers. Project Genie, opened to Google AI Ultra subscribers in January 2026 and worldwide by May, is built on Genie 3, an autoregressive world model generating navigable interactive environments in real time at 720p/24fps (Google DeepMind has not disclosed its parameter count; the frequently cited 11-billion figure belongs to the original 2024 Genie research model). At I/O 2026 Google connected Street View to Genie, enabling simulation of real-world streets.

The strategic purpose is not entertainment but **training infrastructure**. DeepMind’s SIMA 2 agent demonstrably improves at tasks inside Genie-generated worlds without new human examples — the arrangement DeepMind’s leadership calls the “Infinite Training Loop” — and Google explicitly positions Genie as a training and evaluation environment for robots and autonomous systems, feeding the Gemini Robotics family of vision-language-action models already collaborating with partners such as Boston Dynamics and Apptikon. The investment translation is a change of market denomination: video generation is a content-market game measured in tens of billions of dollars; world models are a bid on the *training infrastructure of physical AI*, a market plausibly measured in trillions. The caveats are equally important and equally load-bearing: Genie 3 sustains only minutes of consistent interaction, SIMA 2 still struggles with long-horizon tasks, and reliable sim-to-real transfer remains unproven. The metric to track — directly analogous to the autonomous task horizon that governed the digital-agent economy — is **simulation consistency time**.

6.2 Glasses and microphones: the interface war moves onto the body

AI-native wearables crossed their commercial threshold in 2025: EssilorLuxottica reported a figure of more than seven million Meta AI glasses sold for 2025 — a sharp step up on the prior year, though the basis is not stated precisely enough to fix the multiple — with demand for the display-equipped Ray-Ban model strong enough that Meta paused international expansion on inventory constraints. On IDC’s category series, 2.25 million display-less smart glasses shipped in Q1 2026, with Meta holding 69.2% of the category, against a full-year 2026 forecast of 13.6

million units and roughly \$5.1 billion in revenue. Two measurement caveats attach and are load-bearing for any use of these figures: the seven-million disclosure is a single-source vendor statement (the manufacturing partner of one product line) rather than an industry series, and such statements do not always distinguish cumulative shipments from in-year sales; the IDC series, for its part, covers the display-less category only and is not additive with display-equipped devices. Samsung has confirmed Android XR glasses for 2026, and Apple is widely expected to enter by 2027.

Google's entry defines the strategic logic of the category. Its first consumer Android XR glasses — co-developed with Samsung, styled by Gentle Monster and Warby Parker, launching fall 2026 with iPhone support confirmed — deliberately ship with *no display*: cameras, microphones, and speakers only, with Gemini handling translation, navigation, identification, summarization, and messaging by voice. The design choice reveals the thesis: rather than wait for the hard problem of all-day AR displays, ship the mature components now — always-worn cameras and microphones as **context-collection devices**, voice as the output channel. The essence of the smart-glasses race is not the display; it is the sensor. Whoever is worn all day acquires the deepest personal context, and context is where the durable moat of the AI economy now lives — subject to a real adoption constraint the category has yet to resolve: privacy, consent, on-device processing, and data-retention norms for always-worn cameras and microphones will shape how much of that context can actually be collected and used. Positioning has already split along familiar lines — Meta social-first, Google productivity-and-ecosystem-first — and the wearable is to physical-life context what the desktop agent was to digital-work context: the harness through which the model reaches the substrate.

6.3 The two-axis synthesis

The full arc of this report resolves into the frame introduced in Section 2.3. On Axis 1 — digital work context — the war is in its harvest phase: text-centric, verification-rich, monetized against labor, with strengths favoring Anthropic and OpenAI, and its economics measured token by token in the companion enterprise report. On Axis 2 — physical life context — the war is in its sowing phase: perception- and latency-centric, hardware-distributed, monetized against attention and commerce today and physical labor via robotics tomorrow, with strengths favoring Google and Meta. Generative media is Axis 2's first monetizable beachhead; world models are its training infrastructure; wearables are its sensor network. The expansion of AI “from desktop to mobile” is therefore not a form-factor shift but a bifurcation of the battlefield — and the two axes are linked, because a winner on Axis 2 gains distribution leverage back over Axis 1. Google currently fields the most visibly integrated two-axis position — consumer distribution, generative video, interactive world models, robotics models, and wearable hardware partnerships within a single ecosystem. It is not the only serious participant on any individual layer, but no competitor presently covers all of them at once — in this report's judgment, among the most under-priced structural facts in the 2026 landscape.

7. Risks and Sensitivities

The monetization cap governs the largest segments. The advertising and creator segments — together the bulk of the track’s compute — are bounded by the advertising-and-subscription pool, and the model’s largest single vulnerability is therefore not technical but economic. If attention monetization per generated minute falls faster than generation costs (a plausible outcome of “slop” saturation), realized volume undershoots even as capacity soars — the Sora consumer episode in miniature, replayed at industry scale. Conversely, if agentic commerce routes more purchasing through generated content, the transaction volume that content mediates grows, and the cap rises with it. But the same shift cuts both ways: it makes generated video necessary in more places while threatening to make what that video does — persuade — beside the point.

Volume coefficients carry the widest error bars in this series. The unit-cost side of the model (GPU-hours per minute, per image, per audio-hour) is informed by observed pricing and published throughput but is not directly measurable from either, and the image and audio coefficients additionally combine technically distinct workloads (TTS and music generation share one coefficient) at coarse resolution. As Section 4.3 states, these are engineering assumptions informed by published throughput and pricing but not identifiable from service prices; the residual risk is one-directional, since continued efficiency gains and open-weight self-hosting would bias the coefficients downward over the forecast, lowering the GPU-hour totals without changing the generated-minute volumes. The volume side — creator adoption rates and delivered-minutes-per-video above all — rests on adoption judgments with order-of-magnitude bands, materially wider than any coefficient in the enterprise token model. The segment allocation is more defensible than the total.

Exploration depth is the model’s most consequential single judgment. The two-pass structure of Section 5.1 makes compute proportional to how many candidates a human is willing to look at, and that number rests on a handful of published production accounts rather than on any systematic survey. The risk runs in both directions. If review tooling matures — automated shortlisting of candidates, models that rank their own outputs — exploration depth rises and the creator and professional segments scale beyond these figures. If instead first-time-right rates improve, as vendors claim and as reference-conditioning workflows suggest, depth falls and the totals compress. A flat path for depth per task sits between those cases; a reader holding a strong view either way can scale the affected segments directly. The professional segment carries a separate and larger uncertainty, since its exploration volume per project rises with the number of shots AI reaches rather than with depth per shot.

The social-acceptance assumption is falsifiable and could fail. The model treats labeling mandates and the anti-slop backlash as transitional friction normalizing through 2027–28. A durable regulatory or platform-policy regime that suppresses AI-generated content distribution — mandatory prominent labeling with demonstrated engagement penalties, for instance — would bind the creator-economy and ad-creative segments directly.

Supply-side measurement asymmetries. Seedance distributes primarily through ByteDance’s Doubao ecosystem without a public Western API at the time of writing, limiting independent verification of usage and pricing claims; leaderboard positions reshuffle monthly; and several capability claims circulate through promotional secondary sources. Wearables data carries

the single-source and cumulative-versus-in-year caveats noted in Section 6.2. Institutional use of any single figure herein should rest on primary-source triangulation rather than any single vendor disclosure, per the caveats flagged in Sections 4.3, 6.2, and this section.

World-model risk is concentrated in two unproven links. The trillion-dollar framing of Section 6.1 requires both extended simulation consistency (currently minutes) and reliable sim-to-real transfer (currently undemonstrated at scale). Until both are shown, world models should be valued as an option, not a discounted cash flow.

Interaction with the enterprise track. The two demand functions are modeled as independent, but they compete for the same accelerator supply at the margin. A media demand surprise on the upside would tighten the inference-capacity market that the enterprise track also draws on — a channel the companion infrastructure volume quantifies.

8. Tracking Indicators

The model is built to be falsified on a short horizon. On the content economy: the per-clip and per-second generation cost curve, where each step down pulls forward the media fleet; video-model leaderboard churn and the share of professional workflows on multi-model aggregators; the AI-generated share of published video minutes on major platforms, the direct read on the insert-to-segment-to-full-generation progression; any published data on candidate-to-delivered ratios in production workflows, which is the direct test of the exploration-depth assumption that sets most of the compute total; and ad-platform disclosures of AI-generated creative share in auction volume, the cleanest confirmation of the machine-verified volume engine. On monetization: effective CPM and revenue-per-generated-minute trends for AI-heavy channels versus baseline, the live read on the cap; and any repeat of the Sora pattern — consumer generation products shut or throttled for burn — which marks the cap binding. On the physical-context frontier: simulation consistency time and documented sim-to-real transfers, the Axis-2 task-horizon analogue; smart-glasses unit shipments and platform shares as Android XR launches and Apple's entry timing resolves; and robotics-partnership commercialization milestones in the Gemini Robotics ecosystem.

A closing synthesis. The enterprise report argued that 2026–2028 would be remembered as the interval in which “which industries adopt AI” gave way to “which verification mechanisms automate next.” This report extends the sequence off the desktop: the advertising auction automated first among media workloads because its verifier was already a machine; the segments verified by human senses follow as cost and quality cross their thresholds; and the frontier beyond content — world models and wearables — is the machinery by which AI acquires the physical context that no document ever contained. Token demand was the shadow the first front cast on the infrastructure market. GPU-hours are the shadow of the second — and the companion infrastructure volume reads both shadows onto silicon together.

References

Artificial Analysis. “Text-to-Video Leaderboard” and “Video Model Comparisons” (Seedance, HappyHorse, Kling, Veo Arena Elo and per-minute pricing). Accessed July 2026. <https://artificialanalysis.ai/video/leaderboard/text-to-video>

BuildMVPFast. “Best Text-to-Video AI Generators July 2026: Top 10 Models Ranked” (Sora shutdown dates; Seedance 2.5, Kling 3.0 Turbo, Gemini Omni Flash pricing). July 2026. <https://www.buildmvpfast.com/articles/best-llms-2026-guide/video-generation-ai>

GMI Cloud. “GPU Cloud Cost for AI Inference at Scale in 2026” (managed-inference video pricing ~\$0.07/clip; H100 on-demand rates). March 2026. <https://www.gmicloud.ai/en/blog/gpu-cloud-cost-ai-inference-at-scale>

AI/ML API. “Best AI Video Generators 2026: Veo 3.1, Kling, Sora 2, Seedance & More Compared.” April 2026. <https://aimlapi.com/blog/best-ai-video-generators-2026-veo-3-1-kling-sora-2-seedance-more-compared>

Amazon seller and marketplace statistics. Marketplace Pulse and Seller Assistant, “Amazon Statistics for Sellers 2026” (~1.6M active sellers; third-party GMV). April 2026. <https://www.sellerassistant.app/blog/amazon-statistics-for-sellers-in-2026-key-insights/>

Atlas Cloud. “Seedance vs Kling vs Sora vs Veo: Head-to-Head Comparison.” May 2026. <https://www.atlascloud.ai/blog/guides/seedance-vs-kling-vs-sora-vs-veo>

Campus Technology. “Google Unveils Android XR Smart Glasses, Powered by Gemini AI.” May 22, 2026. <https://campustechnology.com/articles/2026/05/22/google-unveils-android-xr-smart-glasses-powered-by-gemini-ai.aspx>

CoderSera / CatDoes / Composio (agentic-war context for OpenAI resource reallocation): “Claude Code vs OpenAI Codex” comparison series. April–June 2026. <https://codersera.com/blog/claude-code-vs-openai-codex-2026/>

DevTk.AI. “AI Video API Pricing 2026.” February 2026. <https://devtk.ai/en/blog/ai-video-generation-pricing-2026/>

The Gadgeteer. “Google AI Smart Glasses Introduced” (EssilorLuxottica 7M+ Meta AI glasses sold in 2025). April 2026. <https://the-gadgeteer.com/2026/04/27/google-android-xr-ai-smart-glasses/>

Forbes (Liu, P.). “How Much OpenAI May Be Spending on AI-Generated Sora Videos” (outside estimate of Sora inference costs and its stated assumptions). November 10, 2025. <https://www.forbes.com/sites/phoebeliu/2025/11/10/openai-spending-ai-generated-sora-videos/>

Google. “Android XR at I/O 2026” (consumer AI glasses: fall 2026 launch, audio-first design, Android and iPhone support). May 2026. <https://blog.google/products-and-platforms/platforms/android/android-xr-io-2026/>

Google Blog. “Project Genie: AI world model now available for Ultra users.” January–February 2026. <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/project-genie/>

Google Blog. “Project Genie Expands” (worldwide Ultra availability; Street View integration). May 2026. <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/project-genie-expands/>

Google DeepMind. “Genie 3: A New Frontier for World Models” (720p/24fps real-time interactive environments; parameter count undisclosed). <https://deepmind.google/blog/genie-3-a-new-frontier-for-world-models/>

IDC. “Smart Glasses Surge: The XR Market Is Rewriting Its Own Rules” (Q1 2026 display-less shipments 2.25M; Meta share 69.2%; 2026 forecast 13.6M units, ~\$5.1B). 2026. <https://www.idc.com/resource-center/blog/smart-glasses-surge-the-xr-market-is-rewriting-its-own-rules/>

Google Blog. “I/O 2026 developer highlights: Antigravity, Gemini API, AI Studio.” May 2026. <https://blog.google/innovation-and-ai/technology/developers-tools/google-io-2026-developer-highlights/>

Google DeepMind. “Gemini Robotics” model pages and blog series (2025–2026). <https://deepmind.google/models/gemini-robotics/>

Google DeepMind. “Gemini Robotics ER-1.6.” April 2026. <https://blog.google/innovation-and-ai/models-and-research/google-deepmind/gemini-robotics-er-1-6/>

Meta Platforms. “Meta’s AI Product News from Connect” (more than 1 million advertisers using generative AI ad tools; 15 million ads generated in one month). September 2024. <https://about.fb.com/news/2024/09/metas-ai-product-news-connect/>

Meta Platforms. “Performance Talks: How Zenith and Farfetch Are Using AI to Drive Results” (more than 10 million advertising businesses). <https://about.fb.com/news/performance-talks-how-zenith-and-farfetch-are-using-ai-to-drive-results/>

Netflix. Second Quarter 2025 Earnings Interview (first generative-AI final footage in a released title). July 2025. <https://ir.netflix.net/investor-news-and-events/investor-events/event-details/2025/Netflix-Second-Quarter-2025-Earnings-Interview-2025--ZzeVaVp2V/default.aspx>

NVIDIA. “NVIDIA Launches Cosmos 3, the Open Frontier Foundation Model for Physical AI.” 2026. <https://nvidianews.nvidia.com/news/nvidia-launches-cosmos-3-the-open-frontier-foundation-model-for-physical-ai>

OpenAI. “What to Know About the Sora Discontinuation” (app discontinued April 26, 2026; API September 24, 2026). 2026. <https://help.openai.com/en/articles/20001152-what-to-know-about-the-sora-discontinuation>

Runway. “Runway Partners with Lionsgate.” September 2024. <https://runwayml.com/news/runway-partners-with-lionsgate>

Lushbinary. “AI Video Generation 2026: Sora 2 vs Veo 3.1 vs Kling 3.0 Compared.” April 2026. <https://lushbinary.com/blog/ai-video-generation-sora-veo-kling-seedance-comparison/>

Memeburn. “Google AI Smart Glasses 2026: Gemini, iPhone Support, Meta Rival.” May 2026. <https://memeburn.com/google-ai-smart-glasses-2026-gemini-iphone-meta/>

Notebookcheck. “Google Android XR smart glasses launching this fall with Gemini AI.” May 2026. <https://www.notebookcheck.net/Google-Android-XR-smart-glasses-launching-this-fall-with-Gemini-AI.1300619.0.html>

Pinggy. “Best Video Generation AI Models in 2026.” May 2026 (leaderboard details as of the May snapshot). https://pinggy.io/blog/best_video_generation_ai_models/

Stack Overflow. “2025 Developer Survey: AI” (84% of professional developers use or plan to use AI tools; 51% daily use). 2025. <https://survey.stackoverflow.co/2025/ai>

Pixflow. “Best AI Video Generator in 2026.” June 2026. <https://pixflow.net/blog/best-ai-video-generator/>

Shopify. “Global Ecommerce Sales Growth Report 2026” (global e-commerce base and SKU context). 2026. <https://www.shopify.com/blog/global-ecommerce-sales>

Teleprompter. “2025 YouTube Statistics: Global Overview and Key Trends” (~500 hours uploaded per minute; ~65M active creators). December 2025. <https://www.teleprompter.com/blog/2025-youtube-statistics>

YouTube Official Blog. “Second Chances on YouTube” (more than 3 million channels in the YouTube Partner Program). <https://blog.youtube/inside-youtube/second-chances-on-youtube/>

TechCrunch. “Google’s Genie world model can now simulate real streets with Street View.” May 19, 2026. <https://techcrunch.com/2026/05/19/googles-genie-world-model-can-now-simulate-real-streets-with-street-view/>

TechTimes. “DeepMind World Models Train Robots in Imagined Worlds: SIMA Practices Inside Genie 3.” June 6, 2026. <https://www.techtimes.com/articles/317932/20260606/deepmind-world-models-train-robots-imagined-worlds-sima-practices-inside-genie-3-model.htm>

UploadVR. “Google Gemini Smart Glasses Launching in Fall to Take on Ray-Ban Meta.” June 2026. <https://www.uploadvr.com/gemini-smart-glasses-are-launching-in-fall-to-take-on-ray-ban-meta/>

Disclaimer: This report is for informational purposes only and does not constitute investment advice. Quantitative estimates herein are order-of-magnitude modeling assumptions, not measured data, and are documented as adjustable inputs in the media sheet of the companion workbook (Agentic AI Token Demand Model workbook: Section C and the Media Model derivation sheet). Vendor-disclosed figures (pricing, unit shipments, leaderboard positions) carry the verification caveats noted in Sections 6–7.

The Intelligence Economy · Report 10
© 2026 Wisdom Hill. Licensed under CC BY-NC-ND 4.0.
Quotation and summary are freely permitted with attribution.

Cite as: Wisdom Hill Research. 2026. “10. The Functional Anatomy of Consumer AI Demand.”
In The Intelligence Economy, No. 10. Wisdom Hill.
<https://wisdomhill.github.io/intelligence-economy/>